

AMBA® ATP Engine

Guide



AMBA ATP Engine

Guide

Copyright © 2020-2024 Arm Limited or its affiliates. All rights reserved.

Release Information

The following changes have been made to this guide:

Change history			
Date	Issue	Confidentiality	Change
15 May 2020	A	Non-Confidential	First release
7 March 2024	B	Non-Confidential	New features added to support Memory System Resource Partitioning and Monitoring (MPAM) modeling

Proprietary Notice

This document is **NON-CONFIDENTIAL** and any use by you is subject to the terms of this notice.

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word "partner" in reference to Arm's customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

If any of the provisions contained in these terms conflict with any of the provisions of any click through or signed written agreement covering this document with Arm, then the click through or signed written agreement prevails over and supersedes the conflicting provisions of these terms. This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm's trademark usage guidelines at <http://www.arm.com/company/policies/trademarks>.

Copyright © 2020-2024 Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.
110 Fulbourn Road, Cambridge, England CB1 9NJ.
LES-PRE-20349

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by Arm and the party that Arm delivered this document to.

Product Status

The information in this document is final, that is for a developed product.

Web Address

<http://www.arm.com>

Contents

AMBA ATP Engine Guide

Preface

About this guide	viii
Intended audience	viii
Using this guide	ix
Conventions	x
Typographic conventions	x
Numbers	x
Additional reading	xi
Arm publications	xi
Other publications	xi
Feedback	xii
Feedback on this guide	xii
Inclusive language commitment	xii

Chapter 01

Introduction

1.1	About the AMBA ATP Engine guide	1-14
1.2	ATP framework	1-15

Chapter 02

Traffic Profiles

2.1	Traffic Profile Manager	2-18
2.1.1	TPM role	2-18
2.1.2	Host interface	2-18
2.2	Traffic Profile classes	2-20
2.2.1	Requester Traffic Profile	2-20
2.2.2	Checker Traffic Profile	2-21
2.2.3	Delay Traffic Profile	2-21
2.2.4	Completer Traffic Profile	2-22
2.3	ATP FIFO	2-23
2.3.1	ATP FIFO update	2-23

	2.3.2	ATP FIFO send	2-23
	2.3.3	ATP FIFO receive	2-24
	2.3.4	ATP FIFO tracker queue	2-24
2.4		Packet descriptor	2-25
	2.4.1	Address Generation	2-25
	2.4.2	Size Generation	2-25
	2.4.3	Random Generator	2-25
	2.4.4	Packet Tagger	2-25
Chapter 03		Events, Logging, and Statistics	
	3.1	ATP events	3-28
	3.1.1	Requester Traffic Profile Events	3-28
	3.1.2	Delay Traffic Profile Events	3-29
	3.1.3	Checker Traffic Profile Events	3-30
	3.1.4	Completer Traffic Profile Events	3-30
	3.2	Kronos	3-31
	3.2.1	Standalone Execution	3-31
	3.2.2	Mixed Mode Execution	3-31
	3.3	Logging	3-32
	3.4	Statistics	3-33
Chapter 04		Adapters	
	4.1	Adapters overview	4-36
	4.2	Adapter Integration API	4-37
	4.3	Adapter Integration Flow	4-39
	4.4	Adapter design guidelines	4-40
	4.5	gem5 adapter	4-41
	4.5.1	Usage	4-41
	4.5.2	Configuration Parameters	4-41
Appendix 05		AMBA ATP Engine Configuration	
	A.1	Configuration file structure	A-44
	A.1.1	Field type specifiers	A-44
	A.1.2	Supported SI and IEC Standard Rate Units	A-44
	A.1.3	Supported Time Units	A-45
	A.2	Configurations	A-46
	A.2.1	Global Configuration	A-46
	A.2.2	Traffic profile	A-46
	A.2.3	Delay Configuration	A-47
	A.2.4	Completer Configuration	A-47
	A.2.5	FIFO configuration	A-48
	A.2.6	Pattern configuration	A-49
	A.2.7	Random descriptor	A-50
Appendix 06		Revisions	

Preface

This preface introduces the *AMBA ATP Engine Guide*. It contains the following sections:

- *About this guide* on page viii
- *Using this guide* on page ix
- *Conventions* on page x
- *Additional reading* on page xi
- *Feedback* on page xii

About this guide

This guide is a software architecture and functional description for the AMBA ATP Engine.

Intended audience

This guide is written for hardware and software engineers who want to understand the AMBA ATP Engine and want to learn how to integrate it into third-party applications.

Using this guide

This guide is organized into the following chapters:

Chapter 1 Introduction

Introduces the AMBA ATP Engine guide.

Chapter 2 Traffic Profiles

Describes the components of the AMBA ATP Engine.

Chapter 3 Events, Logging, and Statistics

Describes ATP Events that can be triggered by a Traffic Profile and propagated by the Traffic Profile Manager to other Profiles which are subscribed to it.

Chapter 4 Adapters

Describes the use of adapters to allow the AMBA ATP Engine to be integrated into a host platform.

Appendix A AMBA ATP Engine Configuration

Defines a standard format for the ATP Engine configuration file. It is structured in a human-readable and machine-readable format.

Appendix B Revisions

Information about the technical changes between released issues of this guide.

Conventions

The following sections describe conventions that this guide uses:

- [Typographic conventions](#)
- [Numbers](#)

Typographic conventions

The typographical conventions are:

<i>italic</i>	Highlights important notes, introduces special terminology, and denotes internal cross-references and citations.
bold	Denotes signal names, and is used for terms in descriptive lists, where appropriate.
monospace	Used for assembler syntax descriptions, pseudocode, and source code examples. Also used in the main text for instruction mnemonics and for references to other items appearing in assembler syntax descriptions, pseudocode, and source code examples.
SMALL CAPITALS	Used for a few terms that have specific technical meanings.

Numbers

Numbers are normally written in decimal. Binary numbers are preceded by 0b, and hexadecimal numbers by 0x. Both are written in a monospace font.

Additional reading

This section lists relevant publications from Arm. See Arm Developer, <https://developer.arm.com/documentation> for access to Arm documentation.

Arm publications

- *AMBA® Adaptive Traffic Profiles Specification* (ARM IHI 0082)
- *AVL AXI XVC Vector Format* (ARM IHI 0030)
- *Arm® Architecture Reference Manual Supplement Memory System Resource Partitioning and Monitoring (MPAM), for A-profile architecture* (ARM DDI 0598D.b)

Other publications

- *The gem5 simulation system*, <http://www.gem5.org>
- *SNOOPy Calendar Queue* (10.1109/WSC.2000.899756)

Feedback

Arm welcomes feedback on its documentation.

Feedback on this guide

If you have any comments or suggestions for additions and improvements, create a ticket at <https://support.developer.arm.com>.

As part of the ticket, include:

- The title, *AMBA® ATP Engine Guide*.
- The number, ARM IHI 0092B.
- The section name to which your comments refer.
- A concise explanation of your comments.

———— **Note** ————

Arm tests PDFs only in Adobe Acrobat and Acrobat Reader, and cannot guarantee the appearance or behavior of any document when viewed with any other PDF reader.

Inclusive language commitment

Arm values inclusive communities. Arm recognizes that we and our industry have used terms that can be offensive.

Arm strives to lead the industry and create change. Previous issues of this document included terms that can be offensive. We have replaced these terms. If you find offensive terms in this document, please contact terms@arm.com.

Chapter 1

Introduction

This chapter introduces the AMBA ATP Engine guide. It contains the following sections:

- [About the AMBA ATP Engine guide on page 1-14](#)
- [ATP framework on page 1-15](#)

1.1 About the AMBA ATP Engine guide

The AMBA ATP Engine is a platform-independent module that generates synthetic traffic, as defined in the ATP specification. It can be plugged into event-driven or time-driven software modeling, simulation, and testing platforms, such as *gem5*, using a simple API mechanism. See [gem5 adapter on page 4-41](#).

The AMBA ATP Engine provides a framework that is designed to generate synthetic traffic in a simple way. It uses standardized configuration files that can be shared across teams and is both human and machine readable.

The AMBA ATP Engine module is platform-independent that can be plugged into a host tool or run as a standalone tool.

1.2 ATP framework

The process of using the ATP framework is:

- Configure .atp files
- Execute the simulation
- Collect the resulting statistics

The AMBA ATP Engine module supports configuration of any number of ATP Requesters and Completer by using external profile files. Profiles elements in these files define the ATP traffic, which uses a FIFO buffer model. A requester can have one or more Profiles. Completers are defined by a single Profile.

Executing the simulation can be done through a host platform or be set up as a standalone program.

Chapter 2

Traffic Profiles

This chapter describes the components of the AMBA ATP Engine. It contains the following sections:

- *Traffic Profile Manager* on page 2-18
- *Traffic Profile classes* on page 2-20
- *ATP FIFO* on page 2-23
- *Packet descriptor* on page 2-25

2.1 Traffic Profile Manager

A Traffic Profile Manager (TPM) is an entity that manages all the configured Traffic Profiles and arbitrates their activation and deactivation. The TPM requests and delivers packets from and to Traffic Profiles. In addition, the TPM relays events to and from subscribing Traffic Profiles.

Figure 2-1 shows an example of a simple TPM embedded in a host system:

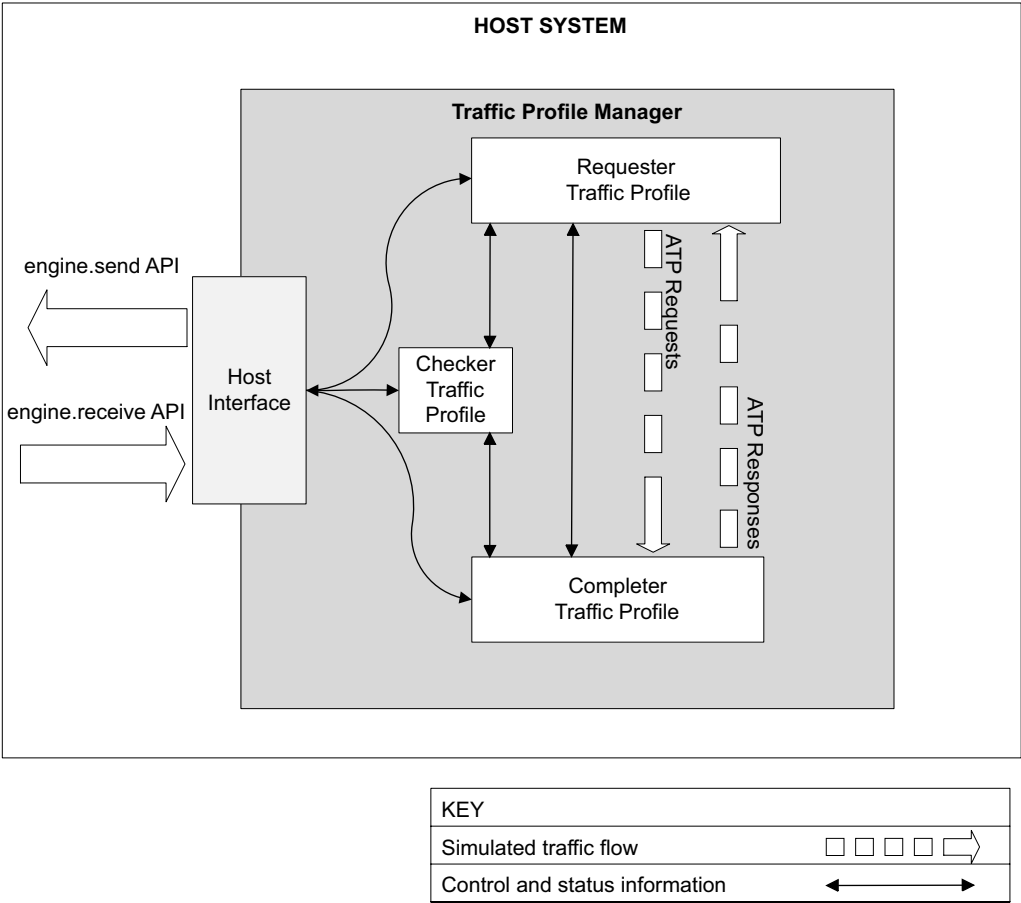


Figure 2-1 Example Traffic Profile Manager simulation

2.1.1 TPM role

Host platforms connect to the TPM in the AMBA ATP Engine. The TPM exposes several APIs that allow host platforms to:

- Send packets
- Receive packets
- Configure specific features
- Retrieve global and per-requester statistics

2.1.2 Host interface

The TPM implements the interfaces that enable the host platform to request and send packets to the AMBA ATP Engine.

The Host Interface defines two APIs to send and receive information.

- send** The send API is called by the host platform to request packets from the AMBA ATP Engine.
- receive** The receive API is used to convey packets to the AMBA ATP Engine, which are then to be delivered to its profiles.

The host interface functions use stream operations. A Stream is defined as a buffered sequence of packets defined by a Traffic Profile on a *First-In First-Out* (FIFO) basis. The packets are chained through wait_for on TERMINATION events.

Table 2-1 shows the host interface operations:

Table 2-1 Traffic Profile Manager operations

Operation	Description
Address Stream Reconfiguration	Allows reconfiguring base and address range of a stream. The parameters for Read and Write operations from the same Stream can be configured independently.
Address Stream Reset	Allows a Stream to be reset to its initial state. The reset can also occur after termination of its profiles. All Traffic Profiles of the Stream are restarted and traffic generation proceeds as if the profiles were in their initial state.
Address Stream Termination query	Interface to query the TPM about the termination of a stream. The interface then returns the state of all the profiles of a stream and whether they have terminated.

2.2 Traffic Profile classes

Traffic Profiles are the entities that generate packets, which are defined by their ATP FIFO configuration. See [ATP FIFO on page 2-23](#).

Traffic Profiles also receive responses to the packets generated. Several kinds of traffic profile entities are implemented in the AMBA ATP Engine:

Requester Traffic Profile

The Requester Traffic Profile implements a traffic profile that sends packet requests and receives packet responses. One or more traffic profiles, each with their corresponding FIFO, can be combined to implement a specific requester behavior. The FIFOs can be in any combination of series and parallel operation.

Checker Traffic Profile

The ATP Checker records all requests and responses that are generated or received by the profile it is monitoring. It stores this information in its own ATP FIFO.

If the profile being monitored has FIFO underrun or overrun detection, it returns an error. The error indicates that the rate that requests and responses are recorded does not match configured FIFO rate.

Delay Traffic Profile

The ATP Delay Traffic Profile implements a delay block. The profile can be used to force a pause between two traffic profiles. Its behavior is to stay active for the configured amount of time, sending no packets and accepting no responses. At end of the configured delay, it terminates.

Completer Traffic Profile

The Completer Traffic Profile implements a fixed latency and bandwidth memory completer. This memory completer receives requests from, and sends responses to, its registered requesters. If all ATP requesters are registered to at least one ATP completer, the AMBA ATP Engine can be run as a standalone executable.

Figure 2-2 shows Traffic Profiles inheritance:




Figure 2-2 Traffic Profiles inheritance

2.2.1 Requester Traffic Profile

Each Requester Traffic Profile contains a Packet Descriptor object, which describes how its ATP Packets are to be generated. Several options for both address and packet size are available:

Address generation options:

- Start from a base value and generate with a configured linear increment
- Start from a base value and generate with two-dimensional increments (increment and stride)
- Random uniform address generation between configured minimum and maximum values
- Normal distribution address generation around a configured mean with a configured standard deviation

Packet size generation options:

- Constant configured value
- Random uniform size generation between configured minimum and maximum values

- Normal distribution size generation around a configured mean with a configured standard deviation

Addresses and packet sizes can also be required to be aligned by enabling the natural alignment feature. It is also possible to specify that either no-alignment or fixed-alignment used. See [Address Generation on page 2-25](#).

A Requester Traffic Profile also includes an ATP FIFO model. The ATP FIFO model shapes the data produced or consumed by the profiles by configuring the FIFO fill and FIFO depletion rate levels. This configuration allows setting a maximum buffered data quantity equal to the configured FIFO maximum size. ATP FIFOs are in [ATP FIFO on page 2-23](#).

A Requester Traffic Profile can be in one of the following states:

Terminated	<p>The profile is terminated, it has depleted all data it was scheduled to send, and has no outstanding transactions in flight which are waiting for response.</p> <p>When a Requester Traffic Profile sets its status to Terminated, it emits an ATP Event. This event signals that its state is Terminated. This terminated status persists until the end of the simulation, unless re-activated (reset) by the TPM. See ATP events on page 3-28.</p>
Locked	<p>The profile is not terminated, and is either:</p> <ul style="list-style-type: none"> • Waiting for responses to unlock its outstanding transactions • Waiting for responses and has terminated all its data scheduled to be sent • Waiting for another profile to terminate • Halted is a condition of complete lock, triggered by linking this profile to another. Linking is accomplished by switching between halted and non-halted status on the other profile being locked or unlocked.
Active	<p>The profile is ready to generate data. The Requester Traffic Profile is active if it is not Locked and not Terminated.</p>

2.2.2 Checker Traffic Profile

A Checker Traffic Profile is configured with an ATP FIFO just like a Requester Traffic Profile. It is also set to monitor one or more Requester Traffic Profiles. When one of its monitored Requester Traffic Profiles sends or receives a Packet, the TPM invokes the send or the receive function on the Checker Traffic Profile with the same packet. The Packet being sent to the Checker Traffic Profile its ATP FIFO status level and Outstanding Transactions (OT) to change.

When the Checker Traffic Profile records an ATP FIFO underrun or overrun, one or more profiles do not match its configured ATP FIFO fill or depletion rate.

A Checker Traffic Profile can be in one of the following states:

Terminated	<p>The profile is terminated if no longer active.</p> <p>When a Checker Traffic Profile changes its status to Terminated, it emits an ATP Event signaling the termination. The Checker Traffic Profile persists in the terminated state until the end of the simulation, unless re-activated (reset) by the TPM. See Chapter 3 Events, Logging, and Statistics.</p>
Active	<p>One or more of the monitored Requester Traffic Profiles is not terminated yet.</p>

2.2.3 Delay Traffic Profile

A Delay Traffic Profile configuration is a single latency value, expressed in seconds or SI unit of time. See [Supported SI and IEC Standard Rate Units on page A-44](#).

On activation, the Delay Traffic Profile initializes a counter and then reports its expiration time at the next transmission time. This report causes the TPM to query the Delay Traffic Profile again at its expiration time. The query from the TPM will then terminate.

The Delay Traffic Profile is normally used to trigger another profile activation by making that profile wait on the Delay Traffic Profile termination event.

A Delay Traffic Profile can be in one of the following states:

- Terminated** The profile configured delay is expired, counting from its first activation time.
When a Delay Traffic Profile switches its state to Terminated, it emits an ATP Event, signaling its termination. The Delay Traffic Profile persists in its terminated state until the end of the simulation, unless re-activated (reset) by the TPM. See [ATP events on page 3-28](#).
- Active** The Delay Traffic Profile is not waiting for any other profile to terminate. The first activation time is recorded and used to compute its termination time as first activation time plus configured delay.

2.2.4 Completer Traffic Profile

Completer Traffic Profiles simulate memory completers that receive requests and send responses to its registered Requester Traffic Profiles. When all ATP Requesters are registered to at least one ATP Completer, the AMBA ATP Engine can be run as a standalone executable.

A Completer Traffic Profile can be configured with:

- Rate** The maximum service rate that the Completer Traffic Profile can serve incoming requests
- Latency** Request-to-response latency
- Max OT** Maximum number of parallel memory accesses supported
- Width** Width of the memory in bytes, which determines how many accesses a request corresponds to
- Requester** ATP Requester or Requesters associated with this memory completer

Completer Traffic Profiles include an ATP READ FIFO model, which is used to regulate the rate at which requests are served. On a request reception, the ATP FIFO is queried using its send API. If the FIFO can accommodate the request, an amount of data equal to the request size is marked as *in-flight*. This marking signals that the request is accepted and queued, waiting to be processed.

When the Completer Traffic Profile is requested to send responses, it calls the receive API on its FIFO for an amount of data corresponding to the sizes of the responses about to be issued. In-flight data in the FIFO is marked as received and can be consumed by the FIFO drain rate. The drain rate allows more requests to be served.

A Completer Traffic Profile can be in one of the following states:

- Locked** A Completer Traffic Profile has reached its maximum OT. A completer that is Locked cannot accept any more requests or the Completer Traffic Profile has no responses queued to be sent
- Active** A Completer Traffic Profile is Active if it is not locked

2.3 ATP FIFO

The ATP FIFO model tracks data either written or read by a Requester or Completer device. The model constrains that data to the size of the FIFO buffer. The constraints are:

- The maximum number of outstanding transactions allowed
- The data consumption or production rate

The ATP FIFO exposes two main APIs:

- send
- receive

When one of the APIs is called, the FIFO also runs an update function. The update function recomputes the FIFO fill level based on its configured rate and elapsed time since last call.

Figure 2-3 shows the flow of the ATP FIFO:




Figure 2-3 ATP FIFO

2.3.1 ATP FIFO update

The Update function computes the FIFO level update, based on the configured rate and the elapsed time since the last update call. The result can be a non-integer number. The Update function uses a carry function to save decimal leftovers. The carry function adds any decimal part to a store and truncates the FIFO level update to an integer value. When the carry value reaches 1, the Update function adds 1 to the current level update and the carry value is reset.

When the FIFO level update is computed, the Update function checks whether an underrun or overrun occurs. The update function checks by comparing the level update with the current FIFO fill level:

- If the ATP FIFO is a READ FIFO, then an underrun occurs when the current fill level is less than the computed level update.
 - If an underrun occurs, the FIFO level is reset to zero.
 - If an underrun does not occur, then the FIFO level update is subtracted (drain) from READ FIFO level.
- If the ATP FIFO is a WRITE FIFO, then an overrun occurs when the current fill level plus the computed level update is greater than the configured maximum fill level.
 - If an overrun occurs, the FIFO level is capped to its configured max level.
 - If an overrun does not occur, then the FIFO level update is added (fill) to WRITE FIFO level.

2.3.2 ATP FIFO send

Each time the ATP FIFO send API is invoked, a FIFO Update is triggered. After the Update, the send API checks whether is possible to accommodate for a specific data request in the FIFO, based on the FIFO type:

Read FIFO The send function checks if it is possible to allocate space in the FIFO to read the requested amount of data. If the current level (including any in-flight data) plus the requested data is less than the configured maximum level, then the allocation is successful.

Write FIFO The send function checks if it is possible to remove data from the FIFO to write the requested amount. If the current level (minus any in-flight data) is greater than the requested amount, then the removal is successful.

The send function also tracks of the outstanding transactions, which is the data that is in-flight the receive function has not been called yet.

In addition to returning the success status of the requested operation, the send function also computes and returns the next time a request for the same amount of data will be possible to be served. This computation is done by considering the current fill level, the maximum configured level, and the FIFO update rate.

2.3.3 ATP FIFO receive

When the receive API is invoked, the Update function is called. The receive function records the acknowledgment of data previously marked as in-flight with a send function call.

The outstanding transaction counter is decremented, and a specified amount of data is removed from the ATP FIFO in-flight data counter. That amount is:

- Added to the ATP FIFO level, if it is ATP READ FIFO
- Removed from the ATP FIFO level, if it is ATP WRITE FIFO

2.3.4 ATP FIFO tracker queue

The ATP FIFO Tracker Queue is an optional feature that is included in the AMBA ATP Engine. Its purpose is to give a latency measure based on how the FIFO-achieved rate differs from the configured drain or fill rate.

Tracker Queue latency is the difference between the successful issue time and the expected issue time. The successful issue time is the amount of time the send API takes to get the amount of data requested. The expected issue time is the amount of time the FIFO is configured to take to consume or produce that amount of data.

Figure 2-4 shows the sequence that the ATP Tracker Queue performs its function:

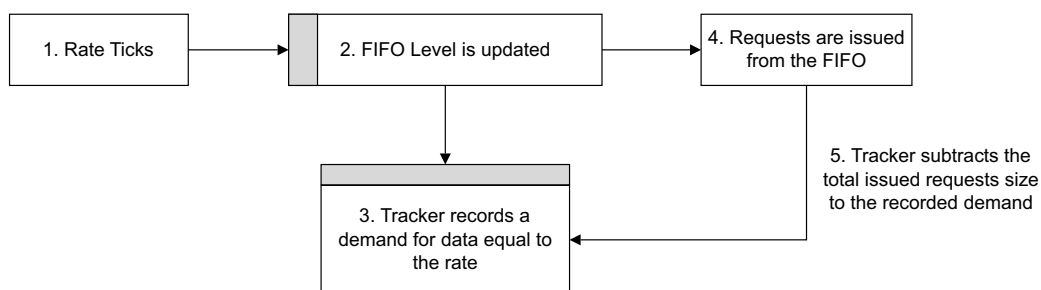


Figure 2-4 ATP Tracker Queue

2.4 Packet descriptor

The Packet Descriptor (PacketDesc class) implements methods to generate packets with specific patterns both for addresses and sizes. When random value generation is required to generate one or both the values, an object of the RandomGenerator class is also used.

The Packet Descriptor also includes a tagger object (PacketTagger) which is used to fill other fields of the generated packets such as local and global IDs.

2.4.1 Address Generation

Address generation can be configured as:

- A linear sequence of addresses with a base plus a fixed increment value (Linear)
- Following one of the patterns available in the Random Generator (Random)

The address generation can be configured with an optional range over which the generated address stream wraps around. In both cases, optional Stride configuration can be provided, which is composed of three parameters:

N	Number of addresses that should be included in a Stride
Increment	Increment to be applied to the addresses of the Stride
Range	Range whose end the Stride addresses that should terminate at

If the Stride parameters are configured, the address generation process:

- Generate one address, using the configured method (Linear or Random).
- Generate N-1 subsequent addresses:
 - Increment the previously generated address by the quantity Increment
 - If the end of the Stride range is reached, then stop generating addresses
 - If the end of the Stride range is not reached, continue generating subsequent addresses

Address alignment can be optionally configured, either to a specific value (must be a power of two), or set to natural.

If a value is provided, generated addresses are masked out so that they are bit-aligned to that value. See [Pattern configuration on page A-49](#).

In natural alignment, a packet address is aligned to the next power of its size.

2.4.2 Size Generation

Size generation of the Packet Descriptor can be configured as a constant value, or as following one of the patterns available in the Random Generator (Random).

2.4.3 Random Generator

The Random Generator class supports two random number distributions, both based on the Mersenne Twister 64-bit generator.

Uniform	Uniformly distributed random numbers that are bounded by the configured min and max values
Normal	Normal distributed (Gaussian) random numbers, with configured average value mean and standard deviation dev

2.4.4 Packet Tagger

The Packet Tagger class assigns values to packet fields that hold metadata other than address and size:

UID	Shared across the entire AMBA ATP Engine. UID is set by the TPM on all outgoing packets and uniquely identifies a packet in the AMBA ATP Engine system, and any connected adapters.
ID	Can be configured with global min/max bounds, or min/max values local to a Traffic Profile

FLOW_ID Used for resource partitioning in gem5; it is used as a placeholder for MPAM Partition ID

IOMMU_ID Used for specifying IOMMU origin in packets

STREAM_ID Used to internally distinguish which requester the packet originated from

Chapter 3

Events, Logging, and Statistics

This chapter describes ATP Events that can be triggered by a Traffic Profile and propagated by the TPM to other Profiles which are subscribed to it. It contains the following sections:

- [ATP events on page 3-28](#)
- [Kronos on page 3-31](#)
- [Logging on page 3-32](#)
- [Statistics on page 3-33](#)

3.1 ATP events

ATP Events can be triggered by a Traffic Profile and propagated by the TPM to other Profiles that are subscribed to the ATP Event. ATP Events signal that something happened to the destination profile at the current ATP time.

Table 3-1 lists all ATP Events with their definitions:

Table 3-1 ATP Events

Event	Definition
NONE	Empty event. Has no effect
ACTIVATION	A profile has become active for the first time
TERMINATION	A profile has issued all the transactions it was configured to issue
FIFO_EMPTY	The FIFO of a profile has become empty
FIFO_FULL	The FIFO of a profile has become full
FIFO_NOT_EMPTY	The FIFO of a profile becomes non-empty after being empty
FIFO_NOT_FULL	The FIFO of a profile becomes non-full after being full
PROFILE_LOCKED	A profile is unable to send data
PROFILE_UNLOCKED	A profile can send data after being unable to
PACKET_REQUEST_RETRY	A rejected request is due to be retried
TICK	Special clock event, triggers TPM tick

All Traffic Profiles derive from the EventManager class, which implements methods to generate and receive events.

When a Traffic Profile subscribes to an Event, the event is added to a wait list in the TPM. When a Traffic Profile generates an Event, it is checked by the TPM against the wait list. The TPM then propagates the Event to all other Traffic Profiles waiting for it.

3.1.1 Requester Traffic Profile Events

Table 3-2 shows the Events that a Requester Traffic Profile can emit:

Table 3-2 Completer Traffic Profile Events

Event	Condition
ACTIVATION	Emitted when: <ul style="list-style-type: none"> The Traffic Profile is created, if it is not configured to wait for any TERMINATION Event. The active API is called, if the API detects that the profile becomes active for the first time.
TERMINATION	Emitted when the active API is called, if it detects that all scheduled data has been sent and there are no outstanding transactions left.
PROFILE_LOCKED	Emitted at the end of a send API call if the Traffic Profile was unable to accommodate the request.
PROFILE_UNLOCKED	Emitted at the end of a send API call if the Traffic Profile could accommodate the request.

Table 3-3 shows the Events that the Requester Traffic Profile can subscribe to:

Table 3-3 Requester Traffic Profile Event subscriptions

Event	Condition
ACTIVATION	A Requester Traffic Profile can wait on: <ul style="list-style-type: none"> One or more other Traffic Profile ACTIVATION events before starting to issue transactions Activation that is externally injected into the Traffic Profile Manager by using its event API
TERMINATION	A Requester Traffic Profile can wait on one or more other Traffic Profile TERMINATION events before starting to issue transactions.
PROFILE_LOCKED	A Requester Traffic Profile can wait on one or more other Traffic Profiles PROFILE_LOCKED events before issuing transactions. The Traffic Profile Requester then starts waiting for a PROFILE_UNLOCKED event and will lock again after the PROFILE_UNLOCKED event occurs.
PROFILE_UNLOCKED	A Requester Traffic Profile can wait on one or more other Traffic Profiles PROFILE_UNLOCKED events before issuing transactions. The Traffic Profile Requester then starts waiting for a PROFILE_LOCKED event and will lock again after the PROFILE_LOCKED event occurs.

3.1.2 Delay Traffic Profile Events

Table 3-4 shows the Events that a Delay Traffic Profile can emit:

Table 3-4 Delay Traffic Profile Events

Event	Condition
ACTIVATION	Emitted when: <ul style="list-style-type: none"> The Traffic Profile is created if it is not configured to wait for any TERMINATION Event When its active API is called and the Traffic Profile is detected to be active for the first time
TERMINATION	Emitted when the Traffic Profile is detected as expired on its active API call. The expiration is defined as when the current time is equal to or past the Traffic Profile deadline. The Traffic Profile deadline is equal to the activation time plus the configured delay.

Table 3-5 shows the Events that the Delay Traffic Profile can subscribe to:

Table 3-5 Delay Traffic Profile Event subscriptions

Event	Condition
ACTIVATION	A Delay Traffic Profile can wait one or more other Traffic Profile ACTIVATION events before starting issuing transactions.
TERMINATION	A Delay Traffic Profile can wait one or more other Traffic Profile TERMINATION events before starting issuing transactions.

3.1.3 Checker Traffic Profile Events

Table 3-6 shows the Events that a Checker Traffic Profile can emit:

Table 3-6 Checker Traffic Profile Events

Event	Condition
TERMINATION	It is emitted when the Traffic Profile is no longer waiting for the TERMINATION of any of its monitored profiles.

Table 3-7 shows the Events that a Checker Traffic Profile can subscribe to:

Table 3-7 Checker Traffic Profile Event subscriptions

Event	Condition
TERMINATION	A Checker Traffic Profile waits on the TERMINATION events of all its monitored profiles.

3.1.4 Completer Traffic Profile Events

Table 3-8 shows the Events that a Completer Traffic Profile can emit:

Table 3-8 Completer Traffic Profile Events

Event	Condition
ACTIVATION	A Completer Traffic Profile emits the activation signal when it is first created.

A Completer Traffic Profile does not subscribe to any events.

3.2 Kronos

The Kronos is a *Future Event List* (FEL) system for the AMBA ATP Engine. It is available to the TPM for scheduling Events and retrieving the next event tick to set the ATP clock to.

Events in Kronos are organized using a calendar queue, which is a fast-lookup data structure specialized in handling timelines.

Kronos enables the AMBA ATP Engine to be run as a standalone executable and in adapter-driven mixed memory mode.

3.2.1 Standalone Execution

When the AMBA ATP Engine is compiled as a standalone executable, it can load ATP files and run them against an internal completer. The standalone executable is named `atpeng`, by default.

This internal completer is configured with:

- No maximum OT limit
- Latency and rate values either taken from:
 - The command line
 - Set to the defaults: 80ns and 32GB/s

The AMBA ATP Engine executable command-line syntax is as follows:

```
./atpeng [.atp file]+ <rate> <latency>
```

If any of the ATP files listed in the command line also contain one or more ATP Completers, those completers will also be instantiated with the default. Any requester already associated with such completers defined in the ATP files will not be associated with the default completer.

The AMBA ATP Engine simulates a system where all requesters are connected to their associated completers with an ideal link. The execution of the simulation ends when the last configured packet is successfully responded to.

3.2.2 Mixed Mode Execution

When ATP is compiled with a host platform, it supports ATP configured internal completers running alongside any host platform defined completer memory systems.

Any requester not associated with internal completers is associated with the host platform by default. All its requests are routed to the platform adapter and incoming responses will be delivered to it.

Mixed Mode Execution is triggered when at least one internal and one external completer are both present.

The AMBA ATP Engine implements an internal routing mechanism to serve packets to internal completers. This mechanism is triggered when its send and receive API are called by the driving adapter. The Mixed Mode execution is transparent to the adapter and no special handling is required.

3.3 Logging

Logging in ATP is delegated to the `Logger` class. The `Logger` follows the singleton design pattern, and as such it provides a `get` API to access and interact with the only instance of the `Logger` class.

The `Logger` class outputs formatted strings to a configured ostream object, which is set to `cout` by default.

The `Logger` defines various logging verbosity levels, of increasing priority. The `Logger` can be configured with one of such levels, causing all log messages logged with inferior priority not to be output to the internal ostream object.

The `Logger` logging levels are, in increasing priority order:

1. `DEBUG_LEVEL`
2. `WARNING_LEVEL`
3. `ERROR_LEVEL`
4. `PRINT_LEVEL`

The `Logger` header file also defines some macro functions which log a message with predefined priority:

1. `LOG`
2. `WARN`
3. `ERROR`, which also causes the AMBA ATP Engine or platform to terminate with `exit(1)`
4. `PRINT`

3.4 Statistics

Traffic Profile Descriptors include a Stats object which logs metrics when data transmission or reception occurs for a profile. The TPM can aggregate Stats objects by a requester on request, by calling its getMasterStats API.

Table 3-9 shows the metrics that are collected for each Traffic Profile:

Table 3-9 Collected Metrics

Name	Description	Unit
Start	Traffic Profile Activation time	Seconds
Time	Traffic Profile Termination time or Statistics collection time, whichever occurred first	Seconds
Sent	Number of ATP Packets sent	-
Received	Number of ATP Packets received	-
Data Sent	Total data sent	Bytes
Data Received	Total data received	Bytes
Latency	Packet Request to Response average latency	Seconds
Jitter	Inter-Request Jitter is computed iteratively, based on RFC 1889 January 1996 $J = J + (D(n-1, n) - J) / 16$	Seconds
Send Rate	Computed as Data Sent/Time	Bytes/Seconds
Receive Rate	Computed as Data Received/Time	Bytes/Seconds
Underruns	Total ATP FIFO underruns	-
Overruns	Total ATP FIFO overruns	-
OT	Average number of outstanding transactions	-
FIFO Level	Average FIFO fill level	Bytes

Chapter 4

Adapters

This chapter describes the use of adapters to allow the AMBA ATP Engine to be integrated into a host platform.

It contains the following sections:

- [*Adapters overview* on page 4-36](#)
- [*Adapter Integration API* on page 4-37](#)
- [*Adapter Integration Flow* on page 4-39](#)
- [*Adapter design guidelines* on page 4-40](#)
- [*gem5 adapter* on page 4-41](#)

4.1 Adapters overview

The AMBA ATP Engine can be integrated into a host platform to use the memory system of the platform. Using this host platform, the ATP Requester is enabled to send memory requests to the host platform and receive memory responses from it.

To integrate the AMBA ATP Engine into the target platform, an adapter layer is required. The adapter is an interface between the platform APIs and the AMBA ATP Engine APIs.

4.2 Adapter Integration API

The integration procedure is described in this section. A gem5 adapter is provided with the official ATP distribution. gem5 adapter integration is detailed in [gem5 adapter on page 4-41](#).

[Table 4-1](#) shows the AMBA ATP Engine APIs that are available to the host platform adapter to set AMBA ATP Engine options:

Table 4-1 Adapter Integration APIs

API function	API	Purpose
Initialization	Load	Loads an ATP file into the AMBA ATP Engine
Control	enableProfilesAsMasters	Special mode where all ATP Profiles are treated as single requesters
	enableTrackerLatency	Enables the ATP Latency Tracker
	setTime	Sets the ATP Time
Stream Management	addressStreamReconfigure	Reconfigures address ranges of a chain of profiles linked by termination waited_for events. The reconfiguration is calculated from the base and range passed, so that they fit the new address space. The calculation starts from the root profile provided.
	streamReset	Resets the stream for a profile. Looks up a list of profiles linked by TERMINATION event and issues them the reset command. <i>Experimental</i>
	streamTerminated	Checks if a profiles stream is terminated. Traverses a profiles stream hierarchy and returns true only if all profiles of the stream have terminated. <i>Experimental</i>
	uniqueStream	Gets a unique reference to a stream. If the requested stream is already in use, this function will clone the stream. <i>Experimental</i>

Table 4-1 Adapter Integration APIs (continued)

API function	API	Purpose
Status	getTimeResolution	Returns the AMBA ATP Engine time resolution See Supported Time Units on page A-45
	getMasters	Returns a set of all configured requester names
	isTerminated	Returns True if all packets for requester <i>m</i> have been transmitted
	waiting	Returns True if the TPM is waiting for packet responses
	getMasterStats	Returns a Stat object for requester containing a snapshot of all its ATP stats computed at the current ATP time
	getMasterSlaves	Returns a map with requester to internal completers pairings
Kronos Control	loop	Run the internal routing loop, which dispatches packets between internal requesters and internal completers
Packet transmission	send	Returns packets generated by active ATP requesters with external (host platform) system destinations.
	receive	Routes packets generated by the host platform to ATP Profiles

4.3 Adapter Integration Flow

To integrate the AMBA ATP Engine into a host platform, it is necessary to follow the code flow:

1. Configure a Traffic Profile Manager object in the adapter by using the control APIs
2. Load ATP input files in the Traffic Profile Manager object using the load API
3. Get the time resolution from the Traffic Profile Manager, which is determined based on the loaded ATP files
4. Start the interaction with the AMBA ATP Engine by calling the send API
5. Route any Packets available to the AMBA ATP Engine by calling the receive API. This call can be triggered asynchronously when a packet to be sent to the AMBA ATP Engine is available.

4.4 Adapter design guidelines

The interactions between the adapter and the AMBA ATP Engine should happen following the guidelines:

- The Traffic Profile Manager object should be embedded into an adapter layer which bridges the AMBA ATP Engine and the host platform.
- The AMBA ATP Engine should exchange information with the adapter layer in the form of Google Protocol Buffer objects. ATP Packets and Stats are Google Protocol Buffer objects and should be exchanged as such.
- Where conversion from ATP Packet format to host packet format is needed, that should be implemented in the adapter layer.
- Any additional information that is required to merge the ATP Packet stream into the host platform, must be resolved at the adapter layer.
- The AMBA ATP Engine relies on the host platform to provide time ticks. Discrete time is acceptable, but non-increasing time ticks are not.
- The host platform should detect whether the AMBA ATP Engine has exhausted all its traffic and it is no longer waiting for packets from the host.

When the TPM does not signal a next transmission time through its send API, the cause could be:

- The TPM could be locked waiting for packets from the host.
- The TPM could have depleted all its scheduled traffic.

Checking the TPM status with the waiting or isTerminated API is needed in this case to ascertain whether AMBA ATP Engine has terminated.

4.5 gem5 adapter

ProfileGen is an adapter for the gem5 simulator that is implemented as a MemObject derived class. ProfileGen encapsulates the TrafficProfileManager and allows gem5 to send and receive memory request and response packets from the AMBA ATP Engine.

ProfileGen connects to other gem5 objects by instantiating a configurable number of requester ports that are dedicated to the individual ATP requesters. ProfileGen sends and receives packets belonging to such requesters.

It also implements a local packet buffer used to store ATP packets when their associated requester ports are unavailable for transmission to the gem5 connected system.

4.5.1 Usage

ProfileGen is to be connected to the system by using its requester ports. The typical way to do so is to directly connect its requester ports to an XBar.

ProfileGen can be used as a reference for integrating the gem5 adapter into gem5 scenarios.

4.5.2 Configuration Parameters

Table 4-2 shows the gem5 configuration parameters:

Table 4-2 gem5 configuration parameters

Parameter	Definition
config_files	Comma-separated list of ATP files
exit_when_done	Flag to enable the simulation to terminate when all ATP requesters have depleted their packets
exit_when_one_master_ends	Flags to enable the simulation to terminate when one ATP requester depletes its packets
trace_atp	Flag to enable tracing ATP-generated packets
trace_atp_file	File name where to store the ATP-generated packets trace
trace_gem	Flag to enable gem5 generated packets in time
trace_gem_file	File name where to store gem5 generated packets trace
trace_m3i	Flag to enable tracing of ATP-generated packets in m3i format [1.4, 3]. M3i traces are generated per requester and named after them
trace_m3i_bus	Bus width to be used to fill the m3i file trace
out_of_range_addresses	Flag to allow the reception of packets with out-of-range addresses from the AMBA ATP Engine
uid_routing	Flag to enable fast routing based on ATP UID packet field
profiles_as_masters	If set, all Traffic Profiles are treated as single requesters
tracker_latency	Flag to enable the ATP Latency Tracker queue

Appendix A

AMBA ATP Engine Configuration

This appendix defines a standard format for the ATP Engine configuration file. It is structured in a human-readable and machine-readable format.

It contains the following sections:

- [Configuration file structure on page A-44](#)
- [Configurations on page A-46](#)

A.1 Configuration file structure

The configuration file syntax follows the Google Protocol Buffer text file format. Each object is enclosed in curly brackets {} and properties of the object are specified with the property: value syntax.

The configuration file is organized as a flat list of Traffic Profiles, preceded by a global configuration section defined in [Global Configuration on page A-46](#).

Each Traffic Profile includes:

- A FIFO Configuration block that specifies the parameters of the AMBA Traffic Profile FIFO
- A Pattern Configuration block that specifies what kind of packets and how the Traffic Profile should generate them

See [FIFO configuration on page A-48](#) and [Pattern configuration on page A-49](#).

A.1.1 Field type specifiers

- Optional** Specifying a value for this field can be omitted
- Repeated** This field can be specified multiple times, it is then processed as array of values.
- Required** Specifying this field value is mandatory

A.1.2 Supported SI and IEC Standard Rate Units

[Table A-1](#) shows the rate unit specifiers that can be used:

Table A-1 Unit specifiers

Specifier	Definition
TB/s	10^{12} Bytes/s
GB/s	10^9 Bytes /s
MB/s	10^6 Bytes /s
kB/s	10^3 Bytes /s
TiB/s	2^{40} Bytes /s
GiB/s	2^{30} Bytes /s
MiB/s	2^{20} Bytes /s
KiB/s	2^{10} Bytes /s
B/s	Bytes /s
Tbit/s	10^{12} bits/s
Gbit/s	10^9 bits /s
Mbit/s	10^6 bits /s
kbit/s	10^3 bits /s
Tibit/s	2^{40} bits /s
Gibit/s	2^{30} bits /s
Mibit/s	2^{20} bits /s
Kibit/s	2^{10} bits /s

A.1.3 Supported Time Units

Table A-2 shows the time units supported in time-based parameters:

Table A-2 Supported Time Units

Unit	Definition
CYCLES	[default]
PS	picoseconds
NS	nanoseconds
US	microseconds
MS	milliseconds
S	seconds

A.2 Configurations

A.2.1 Global Configuration

Table A-3 shows the global configuration parameters can be used to change the AMBA ATP Engine behavior:

Table A-3 Global configuration

Parameter	Type	Requirement	Definition
lowId	uint64	Optional	Packet ID low ID value: all packets contain an ID with this value lower bound
highId	uint64	Optional	Packet ID high ID value: all packets contain an ID with this value upper bound
tracing	bool	Optional	Enables ATP packet tracing to file
trace_dir	string	Optional	Specifies the directory name to store ATP packet traces

A.2.2 Traffic profile

Traffic Profiles are identified in the configuration file by the profile tag. Table A-4 shows the fields each profile has:

Table A-4 Traffic Profile

Parameter	Type	Requirement	Definition
type	Type	Optional	Profile type, which configures the profile as type 'READ' or 'WRITE'. This affects the FIFO behavior.
master_id	string	Required	Requester identifier. Multiple profiles can share a requester ID.
fifo	FifoConfiguration	Optional	Configure the FIFO of a Traffic Profile
pattern	PatternConfiguration	Optional	Packets to be generated
delay	DelayConfiguration	Optional	Its presence indicates that the profile being configured is a Delay type
slave	SlaveConfiguration	Optional	Its presence indicates that the profile being configured is a Completer type
wait_for	string	Optional, repeated	Profiles that should complete before this one can activate or events related to the specified profile
name	string	Optional	Profile identifier. If it is not set, it will be generated as profile<#>
iommu_id	uint32	Optional	IOMMU ID to be set for packets generated by this profile
flow_id	uint64	Optional	FlowID to be set for packets generated by this profile, the FlowID is similar in intent to MPAM Partition ID

A.2.3 Delay Configuration

Table A-5 shows Delay Configuration parameters:

Table A-5 Delay Configuration

Parameter	Type	Requirement	Definition
time	string	Required	Can be a floating-point value and can be expressed in any of the supported time units. See Table A-2 on page A-45 .

A.2.4 Completer Configuration

Table A-6 shows completer configuration parameters:

Table A-6 Completer Configuration

Parameter	Type	Requirement	Definition
Rate	string	Required	Maximum rate the completer can process data at. Can be a floating-point value and can be expressed in any of the supported rate units. See Supported SI and IEC Standard Rate Units on page A-44 .
latency	string	Optional	Static latency with which the completer will respond to requests. Can be a floating-point value and can be expressed in any of the supported time units. See Supported Time Units on page A-45 .
random_latency	RandomDesc	Optional	If present, response latencies are randomly generated according to RandomDesc configuration
random_latency_unit	string	Optional	If random latencies are configured, determines their time unit: can be expressed in any of the supported time units. See Supported Time Units on page A-45 . Default to ns.
TxnLimit	uint64	Optional	Limit to the number of outstanding transactions that the completer can accept. Defaults to 1.
TxnSize	uint64	Required	Granularity with which transactions are processed, that is, size of the completer data unit
master	string	Repeated	Requester assigned to this completer

A.2.5 FIFO configuration

Table A-7 shows FIFO configuration parameters:

Table A-7 FIFO configuration

Parameter	Type	Units	Requirement	Definition
Start	StartupLevel	Enum	Optional	The starting level of the FIFO. Defaults to EMPTY for READ FIFO and FULL for WRITE FIFO.
Start	StartupLevel	Enum		
Full	uint64	Bytes	Required	The maximum number of bytes that the FIFO can contain. A value of 0 means unbounded number of bytes.
TxnLimit	uint64	Integer	Optional	This is the limit to the number of outstanding transactions that the profile can generate. Defaults to 1. A value of 0 means unbounded number of outstanding transactions.
total_txn	uint64		Optional	This is the total number of transactions this profile will generate before deactivating itself. A value of 0 for total_txn indicates unbounded number of transactions.
Rate	uint64	Integer	Required	Fill or Depletion rate of this profile. Can be expressed: <ul style="list-style-type: none"> • In units of ATP time • As a rate with unit specifiers compliant to SI or IEC standards See Supported SI and IEC Standard Rate Units on page A-44.
FrameSize	uint64	Bytes	Optional	Specifies the number of bytes transferred by the FIFO.
FrameTime	uint64	Time	Optional	Specifies the amount of time the FIFO will be active for.

A.2.6 Pattern configuration

Table A-8 shows address pattern configuration parameters:

Table A-8 Pattern configuration

Parameter	Type	Requirement	Definition
cmd	PacketCommand	Optional	Packet command type
address	Address	Optional	Address is configured as:
			base Address generation is sequential, starting from this value.
			increment The value address is increased by this value.
			start If start is specified, when address exceeds Xrange, then generation restarts from start, instead of base.
stride	Stride	Optional	Stride is configured as:
			Stride Value that address increments.
			Xrange Maximum allowed address value
random_address	RandomDesc	Optional	If present, addresses are randomly generated according to RandomDesc configuration. Random addresses can be also configured by specifying <ul style="list-style-type: none"> • Random address generation type • Base address • Range
size	uint64	Optional	If present, all generated packets will have this configured size
random_size	RandomDesc	Optional	If present, packet sizes are randomly generated according to RandomDesc configuration
alignment	uint64	Optional	Address alignment to specified size in bytes when addresses are randomly generated. By default, the alignment is “natural”, that is, addresses are aligned to generated packet sizes. alignment has no effect for non-random address sequences.
lowId	uint64	Optional	Packet ID low ID value: all packets generated by the FIFO will contain an ID with this value lower bound. lowId supersedes the global configuration.
highId	uint64	Optional	Packet ID high ID value: all packets generated by the FIFO will contain an ID with this value upper bound. highId supersedes the global configuration

A.2.7 Random descriptor

Table A-9 shows the parameters used for Random descriptors:

Table A-9 Random descriptor			
Parameter	Type	Requirement	Definition
type	Type	Required	Random distribution type, supported Types are: <ul style="list-style-type: none">• UNIFORM• NORMAL• POISSON• WEIBULL
uniform_desc	UniformDesc	Optional	Uniform distribution configuration, takes min possible value and max possible value parameters
normal_desc	NormalDesc	Optional	Normal distribution configuration, takes mean and std_dev parameters
poisson_desc	PoissonDesc	Optional	Poisson distribution configuration, takes mean parameter
weibull_desc	WeibullDesc	Optional	Weibull distribution configuration, takes scale and shape parameters

Appendix B

Revisions

This appendix describes the technical changes between released issues of this guide.

Table B-1 Differences between Issue B Issue A

Change	Location
Addition of new Packet Tagger assign values	Packet Tagger on page 2-25
Addition of new parameters in <i>Traffic profile</i> table	Traffic profile on page A-46

