

AMBA[®] ATP Engine

Guide



AMBA ATP Engine

Guide

Copyright © 2020 Arm Limited or its affiliates. All rights reserved.

Release Information

The following changes have been made to this guide:

| Change history | | | |
|----------------|-------|------------------|---------------|
| Date | Issue | Confidentiality | Change |
| 15 May 2020 | A | Non-Confidential | First release |

Proprietary Notice

This document is **NON-CONFIDENTIAL** and any use by you is subject to the terms of this notice.

Non-Confidential Proprietary Notice

This document is protected by copyright and other related rights and the practice or implementation of the information contained in this document may be protected by one or more patents or pending patent applications. No part of this document may be reproduced in any form by any means without the express prior written permission of Arm. No license, express or implied, by estoppel or otherwise to any intellectual property rights is granted by this document unless specifically stated.

Your access to the information in this document is conditional upon your acceptance that you will not use or permit others to use the information for the purposes of determining whether implementations infringe any third party patents.

THIS DOCUMENT IS PROVIDED “AS IS”. ARM PROVIDES NO REPRESENTATIONS AND NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, INCLUDING, WITHOUT LIMITATION, THE IMPLIED WARRANTIES OF MERCHANTABILITY, SATISFACTORY QUALITY, NON-INFRINGEMENT OR FITNESS FOR A PARTICULAR PURPOSE WITH RESPECT TO THE DOCUMENT. For the avoidance of doubt, Arm makes no representation with respect to, and has undertaken no analysis to identify or understand the scope and content of, patents, copyrights, trade secrets, or other rights.

This document may include technical inaccuracies or typographical errors.

TO THE EXTENT NOT PROHIBITED BY LAW, IN NO EVENT WILL ARM BE LIABLE FOR ANY DAMAGES, INCLUDING WITHOUT LIMITATION ANY DIRECT, INDIRECT, SPECIAL, INCIDENTAL, PUNITIVE, OR CONSEQUENTIAL DAMAGES, HOWEVER CAUSED AND REGARDLESS OF THE THEORY OF LIABILITY, ARISING OUT OF ANY USE OF THIS DOCUMENT, EVEN IF ARM HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

This document consists solely of commercial items. You shall be responsible for ensuring that any use, duplication or disclosure of this document complies fully with any relevant export laws and regulations to assure that this document or any portion thereof is not exported, directly or indirectly, in violation of such export laws. Use of the word “partner” in reference to Arm’s customers is not intended to create or refer to any partnership relationship with any other company. Arm may make changes to this document at any time and without notice.

If any of the provisions contained in these terms conflict with any of the provisions of any click through or signed written agreement covering this document with Arm, then the click through or signed written agreement prevails over and supersedes the conflicting provisions of these terms. This document may be translated into other languages for convenience, and you agree that if there is any conflict between the English version of this document and any translation, the terms of the English version of the Agreement shall prevail.

The Arm corporate logo and words marked with ® or ™ are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. Other brands and names mentioned in this document may be the trademarks of their respective owners. Please follow Arm’s trademark usage guidelines at <http://www.arm.com/company/policies/trademarks>.

Copyright © 2020, Arm Limited (or its affiliates). All rights reserved.

Arm Limited. Company 02557590 registered in England.
110 Fulbourn Road, Cambridge, England CB1 9NJ.
LES-PRE-20349

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by ARM and the party that ARM delivered this document to.

Product Status

The information in this document is final, that is for a developed product.

Web Address

<http://www.arm.com>

Contents

AMBA ATP Engine Guide

Preface

| | |
|-------------------------------|------|
| About this guide | viii |
| Intended audience | viii |
| Using this guide | ix |
| Conventions | x |
| Typographic conventions | x |
| Numbers | x |
| Additional reading | xi |
| Arm publications | xi |
| Other publications | xi |
| Feedback | xii |
| Feedback on this guide | xii |

Chapter 01

Introduction

| | |
|---------------------------------------|------|
| 1.1 About AMBA ATP Engine guide | 1-14 |
| 1.2 ATP framework | 1-15 |

Chapter 02

Traffic Profiles

| | |
|-------------------------------------|------|
| 2.1 Traffic Profile Manager | 2-18 |
| 2.1.1 TPM role | 2-18 |
| 2.1.2 Host interface | 2-18 |
| 2.2 Traffic Profile classes | 2-20 |
| 2.2.1 Master Traffic Profile | 2-20 |
| 2.2.2 Checker Traffic Profile | 2-21 |
| 2.2.3 Delay Traffic Profile | 2-21 |
| 2.2.4 Slave Traffic Profile | 2-22 |
| 2.3 ATP FIFO | 2-23 |
| 2.3.1 ATP FIFO update | 2-23 |
| 2.3.2 ATP FIFO send | 2-23 |

| | | | |
|--------------------|-------|--|------|
| | 2.3.3 | ATP FIFO receive | 2-24 |
| | 2.3.4 | ATP FIFO tracker queue | 2-24 |
| 2.4 | | Packet descriptor | 2-25 |
| | 2.4.1 | Address Generation | 2-25 |
| | 2.4.2 | Size Generation | 2-25 |
| | 2.4.3 | Random Generator | 2-25 |
| | 2.4.4 | Packet Tagger | 2-25 |
| Chapter 03 | | Events, Logging, and Statistics | |
| | 3.1 | ATP events | 3-28 |
| | 3.1.1 | Master Traffic Profile Events | 3-28 |
| | 3.1.2 | Delay Traffic Profile Events | 3-29 |
| | 3.1.3 | Checker Traffic Profile Events | 3-30 |
| | 3.1.4 | Slave Traffic Profile Events | 3-30 |
| | 3.2 | Kronos | 3-31 |
| | 3.2.1 | Standalone Execution | 3-31 |
| | 3.2.2 | Mixed Mode Execution | 3-31 |
| | 3.3 | Logging | 3-32 |
| | 3.4 | Statistics | 3-33 |
| Chapter 04 | | Adapters | |
| | 4.1 | Adapters overview | 4-36 |
| | 4.2 | Adapter Integration API | 4-37 |
| | 4.3 | Adapter Integration Flow | 4-39 |
| | 4.4 | Adapter design guidelines | 4-40 |
| | 4.5 | gem5 adapter | 4-41 |
| | 4.5.1 | Usage | 4-41 |
| | 4.5.2 | Configuration Parameters | 4-41 |
| Appendix 05 | | AMBA ATP Engine Configuration | |
| | A.1 | Configuration file structure | A-44 |
| | A.1.1 | Field type specifiers | A-44 |
| | A.1.2 | Supported SI and IEC Standard Rate Units | A-44 |
| | A.1.3 | Supported Time Units | A-45 |
| | A.2 | Configurations | A-46 |
| | A.2.1 | Global Configuration | A-46 |
| | A.2.2 | Traffic profile | A-46 |
| | A.2.3 | Delay Configuration | A-46 |
| | A.2.4 | Slave Configuration | A-47 |
| | A.2.5 | FIFO configuration | A-47 |
| | A.2.6 | Pattern configuration | A-48 |
| | A.2.7 | Random descriptor | A-49 |

Preface

This preface introduces the *AMBA ATP Engine Guide*. It contains the following sections:

- *About this guide* on page viii
- *Using this guide* on page ix
- *Conventions* on page x
- *Additional reading* on page xi
- *Feedback* on page xii

About this guide

Software architecture and functional description for the AMBA ATP Engine.

Intended audience

This guide is written for hardware and software engineers who want to understand the inner workings of the AMBA ATP Engine and learn how to integrate it into third-party applications.

Using this guide

Chapter 1 *Introduction* Introduces the AMBA ATP Engine guide.

Chapter 2 *Traffic Profiles* Details the components of the AMBA ATP Engine.

Chapter 3 *Events, Logging, and Statistics* Describes ATP Events that can be triggered by a Traffic Profile and propagated by the Traffic Profile Manager to other Profiles which are subscribed to it.

Chapter 4 *Adapters* Describes the use of adapters to allow the AMBA ATP Engine to be integrated into a host platform.

Appendix A *AMBA ATP Engine Configuration* Defines a standard format for the ATP Engine configuration file. It is structured in a human-readable and machine-readable format.

Conventions

The following sections describe conventions that this guide uses:

- [Typographic conventions](#)
- [Numbers](#)

Typographic conventions

The typographical conventions are:

| | |
|------------------------|---|
| <i>italic</i> | Highlights important notes, introduces special terminology, and denotes internal cross-references and citations. |
| bold | Denotes signal names, and is used for terms in descriptive lists, where appropriate. |
| <code>monospace</code> | Used for assembler syntax descriptions, pseudocode, and source code examples. Also used in the main text for instruction mnemonics and for references to other items appearing in assembler syntax descriptions, pseudocode, and source code examples. |
| SMALL CAPITALS | Used for a few terms that have specific technical meanings. |

Numbers

Numbers are normally written in decimal. Binary numbers are preceded by `0b`, and hexadecimal numbers by `0x`. Both are written in a monospace font.

Additional reading

This section lists relevant publications.

Arm publications

See the Infocenter <http://infocenter.arm.com>, for access to Arm documentation.

- *AMBA Adaptive Traffic Profiles Specification* (ARM IHI 0082)
- *AVL AXI XVC Vector Format* (ARM IHI 0030)
- *SNOOPy Calendar Queue* (10.1109/WSC.2000.899756)

Other publications

- *The gem5 simulation system*, <http://www.gem5.org>

Feedback

Arm welcomes feedback on its documentation.

Feedback on this guide

If you have comments on the content of this guide, send e-mail to errata@arm.com. Give:

- The title, AMBA ATP Engine Guide
- The number, ARM IHI 0092A
- The page number(s) that your comments apply
- A concise explanation of your comments.

Arm also welcomes general suggestions for additions and improvements.

Chapter 1

Introduction

This chapter introduces the AMBA ATP Engine guide:

- [About AMBA ATP Engine guide on page 1-14](#)
- [ATP framework on page 1-15](#)

1.1 About AMBA ATP Engine guide

The AMBA ATP Engine is a platform-independent module that generates synthetic traffic, as defined in the ATP specification. It can be plugged into event-driven or time-driven software modeling, simulation, and testing platforms, such as *gem5*, using a simple API mechanism. See [gem5 adapter on page 4-41](#).

The AMBA ATP Engine provides a framework that is designed to generate synthetic traffic in a simple way. It uses standardized configuration files that can be shared across teams and is both human and machine readable.

The AMBA ATP Engine module is platform-independent that can be plugged into a host tool or run as a standalone tool.

1.2 ATP framework

The process of using the ATP framework is:

- Configure .atp files
- Execute the simulation
- Collect the resulting statistics

The AMBA ATP Engine module supports configuration of any number of ATP masters and slaves by using external profile files. Profiles elements in these files define the ATP traffic, which uses a FIFO buffer model. A master can have one or more Profiles. Slaves are defined by a single Profile.

Executing the simulation can be done through a host platform or be set up as a standalone program.

Chapter 2

Traffic Profiles

This chapter describes the components of the AMBA ATP Engine:

- *Traffic Profile Manager* on page 2-18
- *Traffic Profile classes* on page 2-20
- *ATP FIFO* on page 2-23
- *Packet descriptor* on page 2-25

2.1 Traffic Profile Manager

A Traffic Profile Manager (TPM) is an entity that manages all the configured Traffic Profiles and arbitrates their activation and deactivation. The TPM requests and delivers packets from and to Traffic Profiles. In addition, the TPM relays events to and from subscribing Traffic Profiles.

Figure 2-1 shows an example of a simple TPM embedded in a host system.

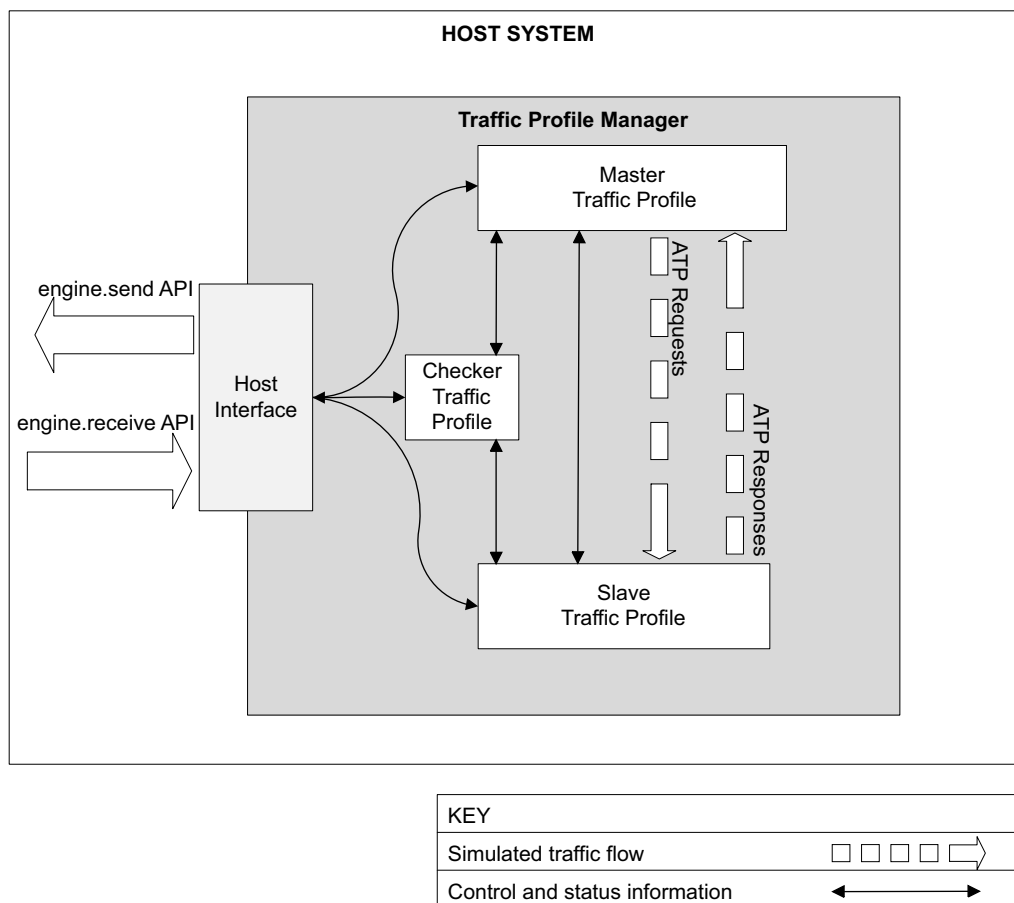


Figure 2-1 Example Traffic Profile Manager simulation

2.1.1 TPM role

Host platforms connect to the TPM in the AMBA ATP Engine. The TPM exposes several APIs that allow host platforms to:

- Send packets
- Receive packets
- Configure specific features
- Retrieve global and per-master statistics

2.1.2 Host interface

The TPM implements the interfaces that enable the host platform to request and send packets to the AMBA ATP Engine.

The Host Interface defines two APIs to send and receive information.

- send** The send API is called by the host platform to request packets from the AMBA ATP Engine.
- receive** The receive API is used to convey packets to the AMBA ATP Engine, which are then to be delivered to its profiles.

The host interface functions use stream operations. A Stream is defined as a buffered sequence of packets defined by a Traffic Profile on a *First-In First-Out* (FIFO) basis. The packets are chained through wait_for on TERMINATION events.

Table 2-1 shows the host interface operations.

Table 2-1 Traffic Profile Manager operations

| Operation | Description |
|----------------------------------|---|
| Address Stream Reconfiguration | Allows reconfiguring base and address range of a stream. The parameters for Read and Write operations from the same Stream can be configured independently. |
| Address Stream Reset | Allows a Stream to be reset to its initial state. The reset can also occur after termination of its profiles. All Traffic Profiles of the Stream are restarted and traffic generation proceeds as if the profiles were in their initial state. |
| Address Stream Termination query | Interface to query the TPM about the termination of a stream. The interface then returns the state of all the profiles of a stream and whether they have terminated. |

2.2 Traffic Profile classes

Traffic Profiles are the entities that generate packets, which are defined by their ATP FIFO configuration. See [ATP FIFO on page 2-23](#).

Traffic Profiles also receive responses to the packets generated. Several kinds of traffic profile entities are implemented in the AMBA ATP Engine:

Master Traffic Profile

The Master Traffic Profile implements a traffic profile that sends packet requests and receives packet responses. One or more traffic profiles, each with their corresponding FIFO, can be combined to implement a specific master behavior. The FIFOs can be in any combination of series and parallel operation.

Checker Traffic Profile

The ATP Checker records all requests and responses that are generated or received by the profile it is monitoring. It stores this information in its own ATP FIFO.

If the profile being monitored has FIFO underrun or overrun detection, it returns an error. The error indicates that the rate that requests and responses are recorded does not match configured FIFO rate.

Delay Traffic Profile

The ATP Delay Traffic Profile implements a delay block. The profile can be used to force a pause between two traffic profiles. Its behavior is to stay active for the configured amount of time, sending no packets and accepting no responses. At end of the configured delay, it terminates.

Slave Traffic Profile

The Slave Traffic Profile implements a fixed latency and bandwidth memory slave. This memory slave receives requests from, and sends responses to, its registered masters. If all ATP masters are registered to at least one ATP slave, the AMBA ATP Engine can be run as a standalone executable.

Figure 2-2 shows Traffic Profiles inheritance.




Figure 2-2 Traffic Profiles inheritance

2.2.1 Master Traffic Profile

Each Master Traffic Profile contains a Packet Descriptor object, which describes how its ATP Packets are to be generated. Several options for both address and packet size are available:

Address generation options:

- Start from a base value and generate with a configured linear increment
- Start from a base value and generate with two-dimensional increments (increment and stride)
- Random uniform address generation between configured minimum and maximum values
- Normal distribution address generation around a configured mean with a configured standard deviation

Packet size generation options:

- Constant configured value
- Random uniform size generation between configured minimum and maximum values
- Normal distribution size generation around a configured mean with a configured standard deviation

Addresses and packet sizes can also be required to be aligned by enabling the natural alignment feature. It is also possible to specify that either no-alignment or fixed-alignment used. See [Address Generation on page 2-25](#).

A Master Traffic Profile also includes an ATP FIFO model. The ATP FIFO model shapes the data produced or consumed by the profiles by configuring the FIFO fill and FIFO depletion rate levels. This configuration allows setting a maximum buffered data quantity equal to the configured FIFO maximum size. ATP FIFOs are in [ATP FIFO on page 2-23](#).

A Master Traffic Profile can be in one of the following states:

| | |
|-------------------|---|
| Terminated | <p>The profile is terminated, it has depleted all data it was scheduled to send, and has no outstanding transactions in flight which are waiting for response.</p> <p>When a Master Traffic Profile sets its status to Terminated, it emits an ATP Event. This event signals that its state is Terminated. This terminated status persists until the end of the simulation, unless re-activated (reset) by the TPM. See ATP events on page 3-28.</p> |
| Locked | <p>The profile is not terminated, and is either:</p> <ul style="list-style-type: none"> • Waiting for responses to unlock its outstanding transactions • Waiting for responses and has terminated all its data scheduled to be sent • Waiting for another profile to terminate • Halted is a condition of complete lock, triggered by linking this profile to another. Linking is accomplished by switching between halted and non-halted status on the other profile being locked or unlocked. |
| Active | <p>The profile is ready to generate data. The Master Traffic Profile is active if it is not Locked and not Terminated.</p> |

2.2.2 Checker Traffic Profile

A Checker Traffic Profile is configured with an ATP FIFO just like a Master Traffic Profile. It is also set to monitor one or more Master Traffic Profiles. When one of its monitored Master Traffic Profiles sends or receives a Packet, the TPM invokes the send or the receive function on the Checker Traffic Profile with the same packet. The Packet being sent to the Checker Traffic Profile its ATP FIFO status level and Outstanding Transactions (OT) to change.

When the Checker Traffic Profile records an ATP FIFO underrun or overrun, one or more profiles do not match its configured ATP FIFO fill or depletion rate.

A Checker Traffic Profile can be in one of the following states:

| | |
|-------------------|---|
| Terminated | <p>The profile is terminated if no longer active.</p> <p>When a Checker Traffic Profile changes its status to Terminated, it emits an ATP Event signaling the termination. The Check Traffic Profile persists in the terminated state until the end of the simulation, unless re-activated (reset) by the TPM. See Chapter 3 Events, Logging, and Statistics.</p> |
| Active | <p>One or more of the monitored Master Traffic Profiles is not terminated yet.</p> |

2.2.3 Delay Traffic Profile

A Delay Traffic Profile configuration is a single latency value, expressed in seconds or SI unit of time. See [Supported SI and IEC Standard Rate Units on page A-44](#).

On activation, the Delay Traffic Profile initializes a counter and then reports its expiration time at the next transmission time. This report causes the TPM to query the Delay Traffic Profile again at its expiration time. The query from the TPM will then terminate.

The Delay Traffic Profile is normally used to trigger another profile activation by making that profile wait on the Delay Traffic Profile termination event.

A Delay Traffic Profile can be in one of the following states:

- Terminated** The profile configured delay is expired, counting from its first activation time
When a Delay Traffic Profile switches its state to Terminated, it emits an ATP Event, signaling its termination. The Delay Traffic Profile persists in its terminated state until the end of the simulation, unless re-activated (reset) by the TPM. See [ATP events on page 3-28](#)
- Active** The Delay Traffic Profile is not waiting for any other profile to terminate. The first activation time is recorded and used to compute its termination time as first activation time plus configured delay

2.2.4 Slave Traffic Profile

Slave Traffic Profiles simulate memory slaves that receive requests and send responses to its registered Master Traffic Profiles. When all ATP Masters are registered to at least one ATP Slave, the AMBA ATP Engine can be run as a standalone executable.

A Slave Traffic Profile can be configured with:

- Rate** The maximum service rate that the Slave Traffic Profile can serve incoming requests
- Latency** Request-to-response latency
- Max OT** Maximum number of parallel memory accesses supported
- Width** Width of the memory in bytes, which determines how many accesses a request corresponds to
- Master** ATP Master or Masters associated with this memory slave

Slave Traffic Profiles include an ATP READ FIFO model, which is used to regulate the rate at which requests are served. On a request reception, the ATP FIFO is queried using its send API. If the FIFO can accommodate the request, an amount of data equal to the request size is marked as "in-flight". This marking signals that the request is accepted and queued, waiting to be processed.

When the Slave Traffic Profile is requested to send responses, it calls the receive API on its FIFO for an amount of data corresponding to the sizes of the responses about to be issued. In-flight data in the FIFO is marked as received and can be consumed by the FIFO drain rate. The drain rate allows more requests to be served.

A Slave Traffic Profile can be in one of the following states:

- Locked** A Slave Traffic Profile has reached its maximum OT. A slave that is Locked cannot accept any more requests or the Slave Traffic Profile has no responses queued to be sent
- Active** A Slave Traffic Profile is Active if it is not locked

2.3 ATP FIFO

The ATP FIFO model tracks data either written or read by a Master or Slave device. The model constrains that data to the size of the FIFO buffer. The constraints are:

- The maximum number of outstanding transactions allowed
- The data consumption or production rate

The ATP FIFO exposes two main APIs: send and receive. When one of those APIs is called, the FIFO also runs an update function. The update function recomputes the FIFO fill level based on its configured rate and elapsed time since last call.

Figure 2-3 shows the flow of the ATP FIFO.



Figure 2-3 ATP FIFO

2.3.1 ATP FIFO update

The Update function computes the FIFO level update, based on the configured rate and the elapsed time since the last update call. The result can be a non-integer number. The Update function uses a "carry" function to save decimal leftovers. The carry function adds any decimal part to a store and truncates the FIFO level update to an integer value. When the carry value reaches 1, the Update function adds 1 to the current level update and the carry value is reset.

When the FIFO level update is computed, the Update function checks whether an underrun or overrun occurs. The update function checks by comparing the level update with the current FIFO fill level:

- If the ATP FIFO is a READ FIFO, then an underrun occurs when the current fill level is less than the computed level update.
 - If an underrun occurs, the FIFO level is reset to zero.
 - If an underrun does not occur, then the FIFO level update is subtracted (drain) from READ FIFO level.
- If the ATP FIFO is a WRITE FIFO, then an overrun occurs when the current fill level plus the computed level update is greater than the configured maximum fill level.
 - If an overrun occurs, the FIFO level is capped to its configured max level.
 - If an overrun does not occur, then the FIFO level update is added (fill) to WRITE FIFO level.

2.3.2 ATP FIFO send

Each time the ATP FIFO send API is invoked, a FIFO Update is triggered. After the Update, the send API checks whether is possible to accommodate for a specific data request in the FIFO, based on the FIFO type:

Read FIFO The send function checks if it is possible to allocate space in the FIFO to read the requested amount of data. If the current level (including any in-flight data) plus the requested data is less than the configured maximum level, then the allocation is successful.

Write FIFO The send function checks if it is possible to remove data from the FIFO to write the requested amount. If the current level (minus any in-flight data) is greater than the requested amount, then the removal is successful.

The send function also tracks of the outstanding transactions, which is the data that is in-flight the receive function has not been called yet.

In addition to returning the success status of the requested operation, the send function also computes and returns the next time a request for the same amount of data will be possible to be served. This computation is done by considering the current fill level, the maximum configured level, and the FIFO update rate.

2.3.3 ATP FIFO receive

When the receive API is invoked, the Update function is called. The receive function records the acknowledgment of data previously marked as in-flight with a send function call.

The outstanding transaction counter is decremented, and a specified amount of data is removed from the ATP FIFO in-flight data counter. That amount is:

- Added to the ATP FIFO level, if it is ATP READ FIFO
- Removed from the ATP FIFO level, if it is ATP WRITE FIFO

2.3.4 ATP FIFO tracker queue

The ATP FIFO Tracker Queue is an optional feature that is included in the AMBA ATP Engine. Its purpose is to give a latency measure based on how the FIFO-achieved rate differs from the configured drain or fill rate.

Tracker Queue latency is the difference between the successful issue time and the expected issue time. The successful issue time is the amount of time the send API takes to get the amount of data requested. The expected issue time is the amount of time the FIFO is configured to take to consume or produce that amount of data.

Figure 2-4 shows the sequence that the ATP Tracker Queue performs its function.

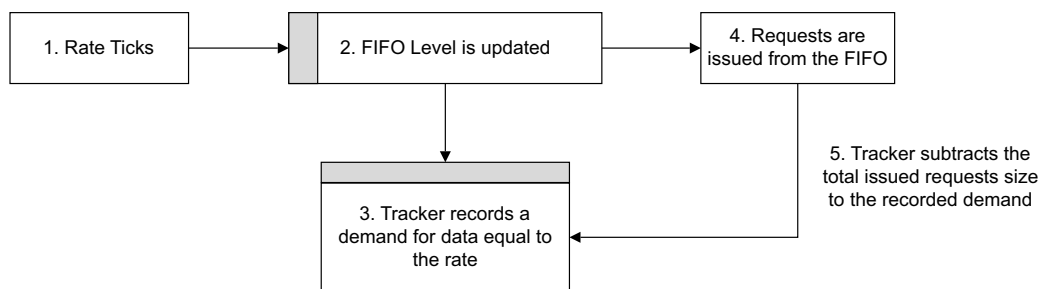


Figure 2-4 ATP Tracker Queue

2.4 Packet descriptor

The Packet Descriptor (PacketDesc class) implements methods to generate packets with specific patterns both for addresses and sizes. When random value generation is required to generate one or both the values, an object of the RandomGenerator class is also used.

The Packet Descriptor also includes a tagger object (PacketTagger) which is used to fill other fields of the generated packets such as local and global IDs.

2.4.1 Address Generation

Address generation can be configured as:

- A linear sequence of addresses with a base plus a fixed increment value (Linear)
- Following one of the patterns available in the Random Generator (Random)

The address generation can be configured with an optional range over which the generated address stream wraps around. In both cases, optional Stride configuration can be provided, which is composed of three parameters:

| | |
|------------------|---|
| N | Number of addresses that should be included in a Stride |
| Increment | Increment to be applied to the addresses of the Stride |
| Range | Range whose end the Stride addresses that should terminate at |

If the Stride parameters are configured, the address generation process:

- Generate one address, using the configured method (Linear or Random).
- Generate N-1 subsequent addresses:
 - Increment the previously generated address by the quantity Increment.
 - If the end of the Stride range is reached, then stop generating addresses
 - If the end of the Stride range is not reached, continue generating subsequent addresses

Address alignment can be optionally configured, either to a specific value (must be a power of two), or set to natural.

If a value is provided, generated addresses are masked out so that they are bit-aligned to that value. See [Pattern configuration on page A-48](#).

In natural alignment, a packet address is aligned to the next power of its size.

2.4.2 Size Generation

Size generation of the Packet Descriptor can be configured as a constant value, or as following one of the patterns available in the Random Generator (Random).

2.4.3 Random Generator

The Random Generator class supports two random number distributions, both based on the Mersenne Twister 64-bit generator.

| | |
|----------------|--|
| Uniform | Uniformly distributed random numbers that are bounded by the configured min and max values |
| Normal | Normal distributed (Gaussian) random numbers, with configured average value mean and standard deviation dev. |

2.4.4 Packet Tagger

The Packet Tagger assigns values to packet fields other than address and size:

| | |
|------------|---|
| UID | Shared across the entire AMBA ATP Engine. UID is set by the TPM on all outgoing packets and uniquely identifies a packet in the AMBA ATP Engine system, and any connected adapters. |
| ID | Can be configured with global min/max bounds, or min/max values local to a Traffic Profile. |

Chapter 3

Events, Logging, and Statistics

This chapter describes ATP Events that can be triggered by a Traffic Profile and propagated by the TPM to other Profiles which are subscribed to it.

- [ATP events on page 3-28](#)
- [Kronos on page 3-31](#)
- [Logging on page 3-32](#)
- [Statistics on page 3-33](#)

3.1 ATP events

ATP Events can be triggered by a Traffic Profile and propagated by the TPM to other Profiles that are subscribed to the ATP Event. ATP Events signal that something happened to the destination profile at the current ATP time.

Table 3-1 shows all ATP Events with their meaning.

Table 3-1 ATP events

| Event | Definition |
|-----------------------------|--|
| NONE | Empty event. Has no effect. |
| ACTIVATION | A profile has become active for the first time |
| TERMINATION | A profile has issued all the transactions it was configured to issue |
| FIFO_EMPTY | The FIFO of a profile has become empty |
| FIFO_FULL | The FIFO of a profile has become full |
| FIFO_NOT_EMPTY | The FIFO of a profile becomes non-empty after being empty |
| FIFO_NOT_FULL | The FIFO of a profile becomes non-full after being full |
| PROFILE_LOCKED | A profile is unable to send data |
| PROFILE_UNLOCKED | A profile can send data after being unable to |
| PACKET_REQUEST_RETRY | A rejected request is due to be retried |
| TICK | Special clock event, triggers TPM tick |

All Traffic Profiles derive from the EventManager class, which implements methods to generate and receive events.

When a Traffic Profile subscribes to an Event, the event is added to a wait list in the TPM. When a Traffic Profile generates an Event, it is checked by the TPM against the wait list. The TPM then propagates the Event to all other Traffic Profiles waiting for it.

3.1.1 Master Traffic Profile Events

Table 3-2 shows the Events that a Master Traffic Profile can emit.

Table 3-2 Slave Traffic Profile Events

| Event | Condition |
|-------------------------|---|
| ACTIVATION | Emitted when: <ul style="list-style-type: none">• The Traffic Profile is created, if it is not configured to wait for any TERMINATION Event,• The active API is called, if the API detects that the profile becomes active for the first time. |
| TERMINATION | Emitted when the active API is called, if it detects that all scheduled data has been sent and there are no outstanding transactions left. |
| PROFILE_LOCKED | Emitted at the end of a send API call if the Traffic Profile was unable to accommodate the request. |
| PROFILE_UNLOCKED | Emitted at the end of a send API call if the Traffic Profile could accommodate the request. |

Table 3-3 shows the Events that the Master Traffic Profile can subscribe to.

Table 3-3 Master Traffic Profile Event subscriptions

| Event | Condition |
|-------------------------|---|
| ACTIVATION | A Master Traffic Profile can wait on: <ul style="list-style-type: none"> One or more other Traffic Profile ACTIVATION events before starting to issue transactions. Activation that is externally injected into the Traffic Profile Manager by using its event API. |
| TERMINATION | A Master Traffic Profile can wait on one or more other Traffic Profile TERMINATION events before starting to issue transactions |
| PROFILE_LOCKED | A Master Traffic Profile can wait on one or more other Traffic Profiles PROFILE_LOCKED events before issuing transactions. The Traffic Profile Master then starts waiting for a PROFILE_UNLOCKED event and will lock again after the PROFILE_UNLOCKED event occurs. |
| PROFILE_UNLOCKED | A Master Traffic Profile can wait on one or more other Traffic Profiles PROFILE_UNLOCKED events before issuing transactions. The Traffic Profile Master then starts waiting for a PROFILE_LOCKED event and will lock again after the PROFILE_LOCKED event occurs. |

3.1.2 Delay Traffic Profile Events

Table 3-4 shows the Events that a Delay Traffic Profile can emit.

Table 3-4 Delay Traffic Profile Events

| Event | Condition |
|--------------------|--|
| ACTIVATION | Emitted when: <ul style="list-style-type: none"> The Traffic Profile is created, if it is not configured to wait for any TERMINATION Event When its active API is called and the Traffic Profile is detected to be active for the first time |
| TERMINATION | Emitted when the Traffic Profile is detected as expired on its active API call. The expiration is defined as when the current time is equal to or past the Traffic Profile deadline. The Traffic Profile deadline is equal to the activation time plus the configured delay. |

Table 3-5 shows the Events that the Delay Traffic Profile can subscribe to.

Table 3-5 Delay Traffic Profile Event subscriptions

| Event | Condition |
|--------------------|---|
| ACTIVATION | A Delay Traffic Profile can wait one or more other Traffic Profile ACTIVATION events before starting issuing transactions. |
| TERMINATION | A Delay Traffic Profile can wait one or more other Traffic Profile TERMINATION events before starting issuing transactions. |

3.1.3 Checker Traffic Profile Events

Table 3-6 shows the Events that a Checker Traffic Profile can emit.

Table 3-6 Checker Traffic Profile Events

| Event | Condition |
|--------------------|---|
| TERMINATION | It is emitted when the Traffic Profile is no longer waiting for the TERMINATION of any of its monitored profiles. |

Table 3-7 shows the Events that a Checker Traffic Profile can subscribe to.

Table 3-7 Checker Traffic Profile Event subscriptions

| Event | Condition |
|--------------------|---|
| TERMINATION | A Checker Traffic Profile waits on the TERMINATION events of all its monitored profiles |

3.1.4 Slave Traffic Profile Events

Table 3-8 shows the Events that a Slave Traffic Profile can emit.

Table 3-8 Slave Traffic Profile Events

| Event | Condition |
|-------------------|---|
| ACTIVATION | A Slave Traffic Profile emits the activation signal when it is first created. |

A Slave Traffic Profile does not subscribe to any events.

3.2 Kronos

The Kronos is a *Future Event List* (FEL) system for the AMBA ATP Engine. It is available to the TPM for scheduling Events and retrieving the next event tick to set the ATP clock to.

Events in Kronos are organized using a calendar queue, which is a fast-lookup data structure specialized in handling timelines.

Kronos enables the AMBA ATP Engine to be run as a standalone executable and in adapter-driven mixed memory mode.

3.2.1 Standalone Execution

When the AMBA ATP Engine is compiled as a standalone executable, it can load ATP files and run them against an internal slave. The standalone executable is named “atpeng”, by default.

This internal slave is configured with:

- No maximum OT limit
- Latency and rate values either taken from:
 - The command line
 - Set to the defaults: 80ns and 32GB/s

The AMBA ATP Engine executable command-line syntax is as follows:

```
./atpeng [.atp file]+ <rate> <latency>
```

If any of the ATP files listed in the command line also contain one or more ATP Slaves, those slaves will also be instantiated with the default. Any master already associated with such slaves defined in the ATP files will not be associated with the default slave.

The AMBA ATP Engine simulates a system where all masters are connected to their associated slaves with an ideal link. The execution of the simulation ends when the last configured packet is successfully responded to.

3.2.2 Mixed Mode Execution

When ATP is compiled with a host platform, it supports ATP configured internal slaves running alongside any host platform defined slave memory systems.

Any master not associated with internal slaves is associated with the host platform by default. All its requests are routed to the platform adapter and incoming responses will be delivered to it.

Mixed Mode Execution is triggered when at least one internal and one external slave are both present.

The AMBA ATP Engine implements an internal routing mechanism to serve packets to internal slaves. This mechanism is triggered when its send and receive API are called by the driving adapter. The Mixed Mode execution is transparent to the adapter and no special handling is required.

3.3 Logging

Logging in ATP is delegated to the `Logger` class. The `Logger` follows the singleton design pattern, and as such it provides a `get` API to access and interact with the only instance of the `Logger` class.

The `Logger` class outputs formatted strings to a configured ostream object, which is set to `cout` by default.

The `Logger` defines various logging verbosity levels, of increasing priority. The `Logger` can be configured with one of such levels, causing all log messages logged with inferior priority not to be output to the internal ostream object.

The `Logger` logging levels are, in increasing priority order:

1. `DEBUG_LEVEL`
2. `WARNING_LEVEL`
3. `ERROR_LEVEL`
4. `PRINT_LEVEL`

The `Logger` header file also defines some macro functions which log a message with predefined priority:

1. `LOG`
2. `WARN`
3. `ERROR`, which also causes the AMBA ATP Engine or platform to terminate with `exit(1)`
4. `PRINT`

3.4 Statistics

Traffic Profile Descriptors include a Stats object which logs metrics when data transmission or reception occurs for a profile. The TPM can aggregate Stats objects by master on request, by calling its getMasterStats API.

Table 3-9 shows the metrics that are collected for each Traffic Profile.

Table 3-9 Collected Metrics

| Name | Description | Unit |
|---------------|---|---------------|
| Start | Traffic Profile Activation time | Seconds |
| Time | Traffic Profile Termination time or Statistics collection time, whichever occurred first. | Seconds |
| Sent | Number of ATP Packets sent | - |
| Received | Number of ATP Packets received | - |
| Data Sent | Total data sent | Bytes |
| Data Received | Total data received | Bytes |
| Latency | Packet Request to Response average latency | Seconds |
| Jitter | Inter-Request Jitter is computed iteratively, based on RFC 1889 January 1996 $J = J + (D(n-1, n) - J) / 16$ | Seconds |
| Send Rate | Computed as Data Sent/Time | Bytes/Seconds |
| Receive Rate | Computed as Data Received/Time | Bytes/Seconds |
| Underruns | Total ATP FIFO underruns | - |
| Overruns | Total ATP FIFO overruns | - |
| OT | Average number of outstanding transactions | - |
| FIFO Level | Average FIFO fill level | Bytes |

Chapter 4

Adapters

This chapter describes the use of adapters to allow the AMBA ATP Engine to be integrated into a host platform:

- [*Adapters overview* on page 4-36](#)
- [*Adapter Integration API* on page 4-37](#)
- [*Adapter Integration Flow* on page 4-39](#)
- [*Adapter design guidelines* on page 4-40](#)
- [*gem5 adapter* on page 4-41](#)

4.1 Adapters overview

The AMBA ATP Engine can be integrated into a host platform to use the memory system of the platform. Using this host platform, the ATP Master is enabled to send memory requests to the host platform and receive memory responses from it.

To integrate the AMBA ATP Engine into the target platform, an adapter layer is required. The adapter is an interface between the platform APIs and the AMBA ATP Engine APIs.

4.2 Adapter Integration API

The integration procedure is described in this section. A gem5 adapter is provided with the official ATP distribution. gem5 adapter integration is detailed in [gem5 adapter on page 4-41](#).

[Table 4-1](#) shows the AMBA ATP Engine APIs that are available to the host platform adapter to set AMBA ATP Engine options.

Table 4-1 Adapter Integration APIs

| API function | API | Purpose |
|-------------------|--------------------------|--|
| Initialization | Load | Loads an ATP file into the AMBA ATP Engine |
| Control | enableProfilesAsMasters | Special mode where all ATP Profiles are treated as single masters |
| | enableTrackerLatency | Enables the ATP Latency Tracker |
| | setTime | Sets the ATP Time |
| Stream Management | addressStreamReconfigure | Reconfigures address ranges of a chain of profiles linked by termination waited_for events. The reconfiguration is calculated from the base and range passed, so that they fit the new address space. The calculation starts from the root profile provided. |
| | streamReset | Resets the stream for a profile. Looks up a list of profiles linked by TERMINATION event and issues them the reset command. <i>Experimental.</i> |
| | streamTerminated | Checks if a profiles stream is terminated. Traverses a profiles stream hierarchy and returns true only if all profiles of the stream have terminated. <i>Experimental.</i> |
| | uniqueStream | Gets a unique reference to a stream. If the requested stream is already in use, this function will clone the stream. <i>Experimental.</i> |

Table 4-1 Adapter Integration APIs (continued)

| API function | API | Purpose |
|---------------------|-------------------|---|
| Status | getTimeResolution | Returns the AMBA ATP Engine time resolution See Supported Time Units on page A-45 |
| | getMasters | Returns a set of all configured master names |
| | isTerminated | Returns True if all packets for master <i>m</i> have been transmitted |
| | waiting | Returns True if the TPM is waiting for packet responses |
| | getMasterStats | Returns a Stat object for master <i>m</i> containing a snapshot of all its ATP stats computed at the current ATP time |
| | getMasterSlaves | Returns a map with master to internal slaves pairings |
| Kronos Control | loop | Run the internal routing loop, which dispatches packets between internal masters and internal slaves |
| Packet transmission | send | Returns packets generated by active ATP masters with external (host platform) system destinations. |
| | receive | Routes packets generated by the host platform to ATP Profiles |

4.3 Adapter Integration Flow

To integrate the AMBA ATP Engine into a host platform, it is necessary to follow the code flow:

1. Configure a Traffic Profile Manager object in the adapter by using the control APIs
2. Load ATP input files in the Traffic Profile Manager object using the load API
3. Get the time resolution from the Traffic Profile Manager, which is determined based on the loaded ATP files
4. Start the interaction with the AMBA ATP Engine by calling the send API
5. Route any Packets available to the AMBA ATP Engine by calling the receive API. This call can be triggered asynchronously when a packet to be sent to the AMBA ATP Engine is available

4.4 Adapter design guidelines

The interactions between the adapter and the AMBA ATP Engine should happen following the guidelines:

- The Traffic Profile Manager object should be embedded into an adapter layer which bridges the AMBA ATP Engine and the host platform.
- The AMBA ATP Engine should exchange information with the adapter layer in the form of Google Protocol Buffer objects. ATP Packets and Stats are Google Protocol Buffer objects and should be exchanged as such.
- Where conversion from ATP Packet format to host packet format is needed, that should be implemented in the adapter layer.
- Any additional information that is required to merge the ATP Packet stream into the host platform, must be resolved at the adapter layer.
- The AMBA ATP Engine relies on the host platform to provide time ticks. Discrete time is acceptable, but non-increasing time ticks are not.
- The host platform should detect whether the AMBA ATP Engine has exhausted all its traffic and it is no longer waiting for packets from the host.

When the TPM does not signal a next transmission time through its send API, the cause could be:

- The TPM could be locked waiting for packets from the host.
- The TPM could have depleted all its scheduled traffic.

Checking the TPM status with the waiting or isTerminated API is needed in this case to ascertain whether AMBA ATP Engine has terminated.

4.5 gem5 adapter

ProfileGen is an adapter for the gem5 simulator that is implemented as a MemObject derived class. ProfileGen encapsulates the Traffic Profile Manager and allows gem5 to send and receive memory request and response packets from the AMBA ATP Engine.

ProfileGen connects to other gem5 objects by instantiating a configurable number of master ports that are dedicated to the individual ATP masters. ProfileGen sends and receives packets belonging to such masters.

It also implements a local packet buffer used to store ATP packets when their associated master ports are unavailable for transmission to the gem5 connected system.

4.5.1 Usage

ProfileGen is to be connected to the system by using its master ports. The typical way to do so is to directly connect its master ports to an XBar.

ProfileGen can be used as a reference for integrating the gem5 adapter into gem5 scenarios.

4.5.2 Configuration Parameters

Table 4-2 gem5 configuration parameters

| Parameter | Definition |
|---------------------------|--|
| config_files | Comma-separated list of ATP files |
| exit_when_done | Flag to enable the simulation to terminate when all ATP masters have depleted their packets |
| exit_when_one_master_ends | Flags to enable the simulation to terminate when one ATP master depletes its packets |
| trace_atp | Flag to enable tracing ATP-generated packets |
| trace_atp_file | File name where to store the ATP-generated packets trace |
| trace_gem | Flag to enable gem5 generated packets in time |
| trace_gem_file | File name where to store gem5 generated packets trace |
| trace_m3i | Flag to enable tracing of ATP-generated packets in m3i format [1.4, 3]. M3i traces are generated per master and named after them |
| trace_m3i_bus | Bus width to be used to fill the m3i file trace |
| out_of_range_addresses | Flag to allow the reception of packets with out-of-range addresses from the AMBA ATP Engine |
| uid_routing | Flag to enable fast routing based on ATP UID packet field |
| profiles_as_masters | If set, all Traffic Profiles are treated as single masters |
| tracker_latency | Flag to enable the ATP Latency Tracker queue |

Appendix A

AMBA ATP Engine Configuration

This appendix defines a standard format for the ATP Engine configuration file. It is structured in a human-readable and machine-readable format.

A.1 Configuration file structure

The configuration file syntax follows the Google Protocol Buffer text file format. Each object is enclosed in curly brackets {} and properties of the object are specified with the property: value syntax.

The configuration file is organized as a flat list of Traffic Profiles, preceded by a global configuration section defined in [Global Configuration on page A-46](#).

Each Traffic Profile includes:

- A FIFO Configuration block that specifies the parameters of the AMBA Traffic Profile FIFO
- A Pattern Configuration block that specifies what kind of packets and how the Traffic Profile should generate them

See [FIFO configuration on page A-47](#) and [Pattern configuration on page A-48](#).

A.1.1 Field type specifiers

- Optional** Specifying a value for this field can be omitted.
- Repeated** This field can be specified multiple times, it is then processed as array of values.
- Required** Specifying this field value is mandatory.

A.1.2 Supported SI and IEC Standard Rate Units

[Table A-1](#) shows the rate unit specifiers that can be used:

Table A-1 Unit specifiers

| Specifier | Definition |
|-----------|-------------------|
| TB/s | 10^{12} Bytes/s |
| GB/s | 10^9 Bytes /s |
| MB/s | 10^6 Bytes /s |
| kB/s | 10^3 Bytes /s |
| TiB/s | 2^{40} Bytes /s |
| GiB/s | 2^{30} Bytes /s |
| MiB/s | 2^{20} Bytes /s |
| KiB/s | 2^{10} Bytes /s |
| B/s | Bytes /s |
| Tbit/s | 10^{12} bits/s |
| Gbit/s | 10^9 bits /s |
| Mbit/s | 10^6 bits /s |
| kbit/s | 10^3 bits /s |
| Tibit/s | 2^{40} bits /s |
| Gibit/s | 2^{30} bits /s |
| Mibit/s | 2^{20} bits /s |
| Kibit/s | 2^{10} bits /s |

A.1.3 Supported Time Units

Table A-2 shows the time units supported in time-based parameters.

Table A-2 Supported Time Units

| Unit | Definition |
|--------|--------------|
| CYCLES | [default] |
| PS | picoseconds |
| NS | nanoseconds |
| US | microseconds |
| MS | milliseconds |
| S | seconds |

A.2 Configurations

A.2.1 Global Configuration

[Table A-3](#) shows the global configuration parameters can be used to change the AMBA ATP Engine behavior.

Table A-3 Global configuration

| Parameter | Type | Requirement | Definition |
|-----------|--------|-------------|--|
| lowId | uint64 | Optional | Packet ID low ID value: all packets contain an ID with this value lower bound |
| highId | uint64 | Optional | Packet ID high ID value: all packets contain an ID with this value upper bound |
| tracing | bool | Optional | Enables ATP packet tracing to file |
| trace_dir | string | Optional | Specifies the directory name to store ATP packet traces |

A.2.2 Traffic profile

Traffic Profiles are identified in the configuration file by the profile tag. [Table A-4](#) shows the fields each profile has.

Table A-4 Traffic Profile

| Parameter | Type | Requirement | Definition |
|-----------|----------------------|--------------------|---|
| type | Type | Optional | Profile type, which configures the profile as type 'READ' or 'WRITE'. This affects the FIFO behavior. |
| master_id | string | Required | Master identifier. Multiple profiles can share a master ID. |
| fifo | FifoConfiguration | Optional | Configure the FIFO of a Traffic Profile |
| pattern | PatternConfiguration | Optional | Packets to be generated |
| delay | DelayConfiguration | Optional | Its presence indicates that the profile being configured is a Delay type |
| slave | SlaveConfiguration | Optional | Its presence indicates that the profile being configured is a Slave type |
| wait_for | string | Optional, repeated | Profiles that should complete before this one can activate or events related to the specified profile |
| name | string | Optional | Profile identifier. If it is not set, it will be generated as profile<#> |

A.2.3 Delay Configuration

[Table A-5](#) shows Delay Configuration parameters.

Table A-5 Delay Configuration

| Parameter | Type | Requirement | Definition |
|-----------|--------|-------------|---|
| time | string | Required | Can be a floating-point value and can be expressed in any of the supported time units. See Table A-2 on page A-45 . |

A.2.4 Slave Configuration

Table A-6 shows slave configuration parameters.

Table A-6 Slave Configuration

| Parameter | Type | Requirement | Definition |
|---------------------|------------|-------------|---|
| Rate | string | Required | Maximum rate the slave can process data at. Can be a floating-point value and can be expressed in any of the supported rate units. See Supported SI and IEC Standard Rate Units on page A-44. |
| latency | string | Optional | Static latency with which the slave will respond to requests. Can be a floating-point value and can be expressed in any of the supported time units. See Supported Time Units on page A-45. |
| random_latency | RandomDesc | Optional | If present, response latencies are randomly generated according to RandomDesc configuration |
| random_latency_unit | string | Optional | If random latencies are configured, determines their time unit: can be expressed in any of the supported time units. See Supported Time Units on page A-45. Default to ns. |
| TxnLimit | uint64 | Optional | Limit to the number of outstanding transactions that the slave can accept. Defaults to 1 |
| TxnSize | uint64 | Required | Granularity with which transactions are processed, that is, size of the slave data unit |
| master | string | Repeated | Masters assigned to this slave |

A.2.5 FIFO configuration

Table A-7 shows FIFO configuration parameters.

Table A-7 FIFO configuration

| Parameter | Type | Units | Requirement | Definition |
|-----------|--------------|---------|-------------|--|
| Start | StartupLevel | Enum | Optional | The starting level of the FIFO. Defaults to EMPTY for READ FIFO and FULL for WRITE FIFO |
| Start | StartupLevel | Enum | | |
| Full | uint64 | Bytes | Required | The maximum number of bytes that the FIFO can contain. A value of 0 means unbounded number of bytes. |
| TxnLimit | uint64 | Integer | Optional | This is the limit to the number of outstanding transactions that the profile can generate. Defaults to 1. A value of 0 means unbounded number of outstanding transactions. |
| total_txn | uint64 | | Optional | This is the total number of transactions this profile will generate before deactivating itself. A value of 0 for total_txn indicates unbounded number of transactions. |

Table A-7 FIFO configuration (continued)

| Parameter | Type | Units | Requirement | Definition |
|-----------|--------|---------|-------------|---|
| Rate | uint64 | Integer | Required | Fill or Depletion rate of this profile. Can be expressed: <ul style="list-style-type: none"> In units of ATP time As a rate with unit specifiers compliant to SI or IEC standards See <i>Supported SI and IEC Standard Rate Units</i> on page A-44. |
| FrameSize | uint64 | Bytes | Optional | Specifies the number of bytes transferred by the FIFO |
| FrameTime | uint64 | Time | Optional | Specifies the amount of time the FIFO will be active for |

A.2.6 Pattern configuration

Table A-8 shows address pattern configuration parameters.

Table A-8 Pattern configuration

| Parameter | Type | Requirement | Definition | |
|----------------|---------------|-------------|---|---|
| cmd | PacketCommand | Optional | Packet command type | |
| address | Address | Optional | Address is configured as: | |
| | | | base | Address generation is sequential, starting from this value. |
| | | | increment | The value address is increased by this value. |
| | | | start | If start is specified, when address exceeds Xrange, then generation restarts from start, instead of base. |
| stride | Stride | Optional | Stride is configured as: | |
| | | | Stride | Value that address increments. |
| | | | Xrange | Maximum allowed address value. |
| random_address | RandomDesc | Optional | If present, addresses are randomly generated according to RandomDesc configuration. Random addresses can be also configured by specifying <ul style="list-style-type: none">Random address generation typeBase addressRange | |
| size | uint64 | Optional | If present, all generated packets will have this configured size | |
| random_size | RandomDesc | Optional | If present, packet sizes are randomly generated according to RandomDesc configuration | |

Table A-8 Pattern configuration (continued)

| Parameter | Type | Requirement | Definition |
|-----------|--------|-------------|---|
| alignment | uint64 | Optional | Address alignment to specified size in bytes when addresses are randomly generated. By default, the alignment is “natural”, that is, addresses are aligned to generated packet sizes. alignment has no effect for non-random address sequences. |
| lowId | uint64 | Optional | Packet ID low ID value: all packets generated by the FIFO will contain an ID with this value lower bound. lowId supersedes the global configuration. |
| highId | uint64 | Optional | Packet ID high ID value: all packets generated by the FIFO will contain an ID with this value upper bound. highId supersedes the global configuration |

A.2.7 Random descriptor

Table A-9 shows the parameters used for Random descriptors.

Table A-9 Random descriptor

| Parameter | Type | Requirement | Definition |
|--------------|-------------|-------------|--|
| type | Type | Required | Random distribution type, supported Types are: <ul style="list-style-type: none"> • UNIFORM • NORMAL • POISSON • WEIBULL |
| uniform_desc | UniformDesc | Optional | Uniform distribution configuration, takes min possible value and max possible value parameters |
| normal_desc | NormalDesc | Optional | Normal distribution configuration, takes mean and std_dev parameters |
| poisson_desc | PoissonDesc | Optional | Poisson distribution configuration, takes mean parameter |
| weibull_desc | WeibullDesc | Optional | Weibull distribution configuration, takes scale and shape parameters |

